

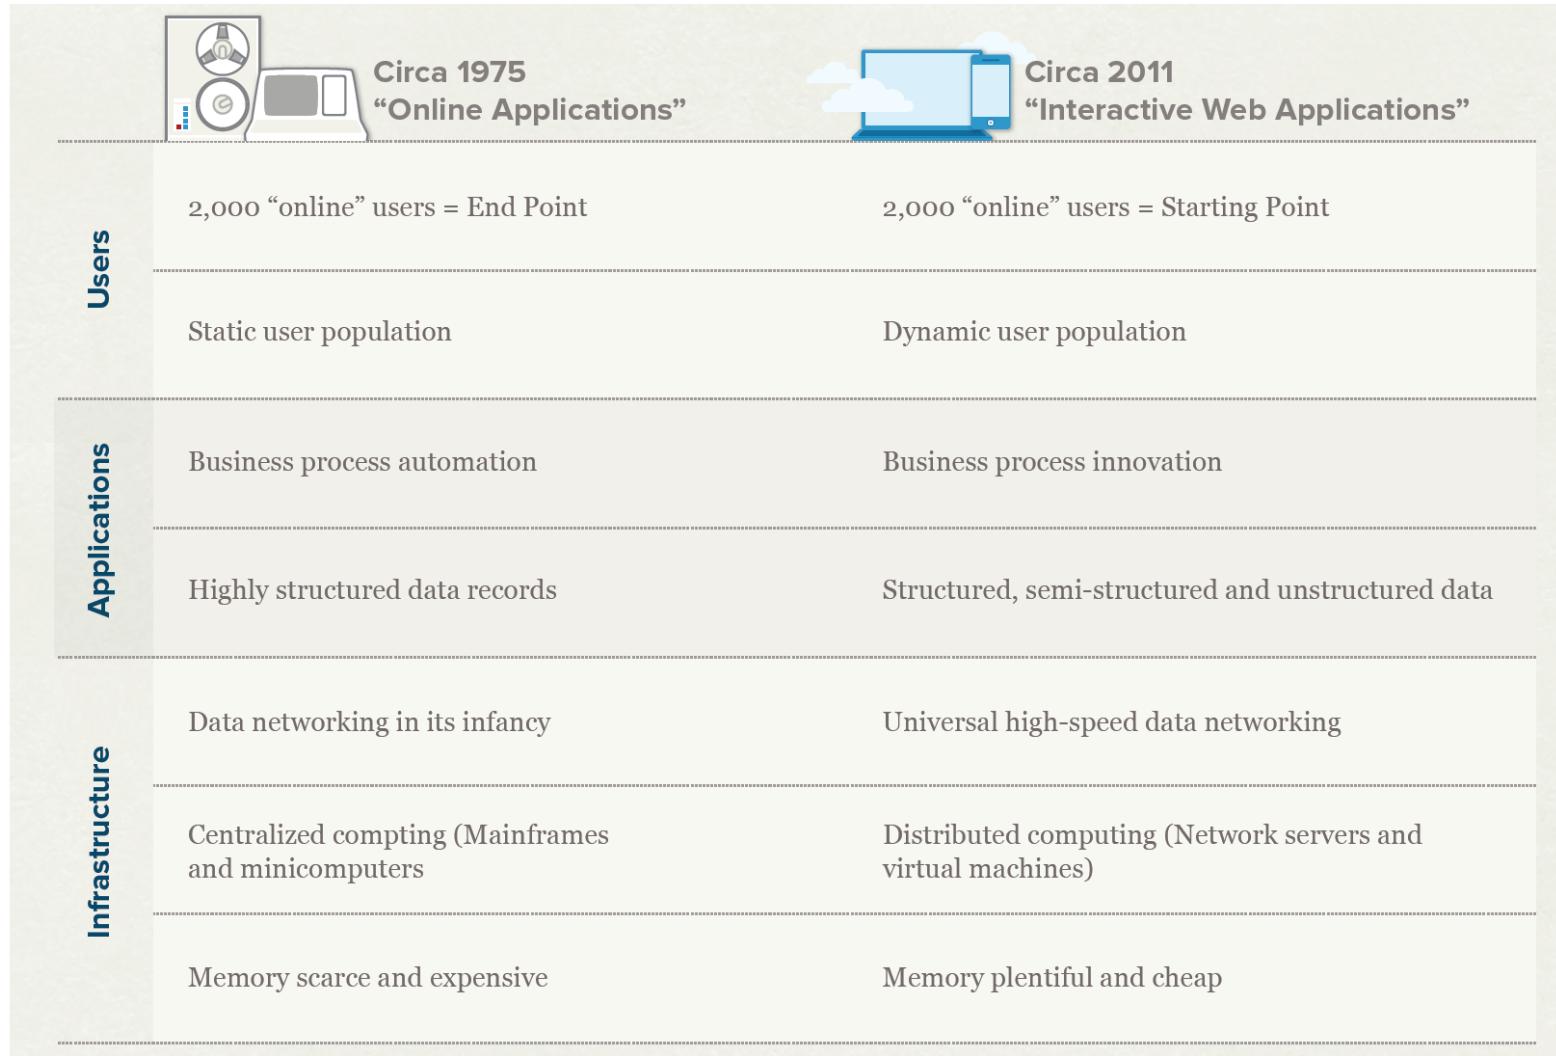
NoSQL database technology: A survey and comparison of systems

James Phillips

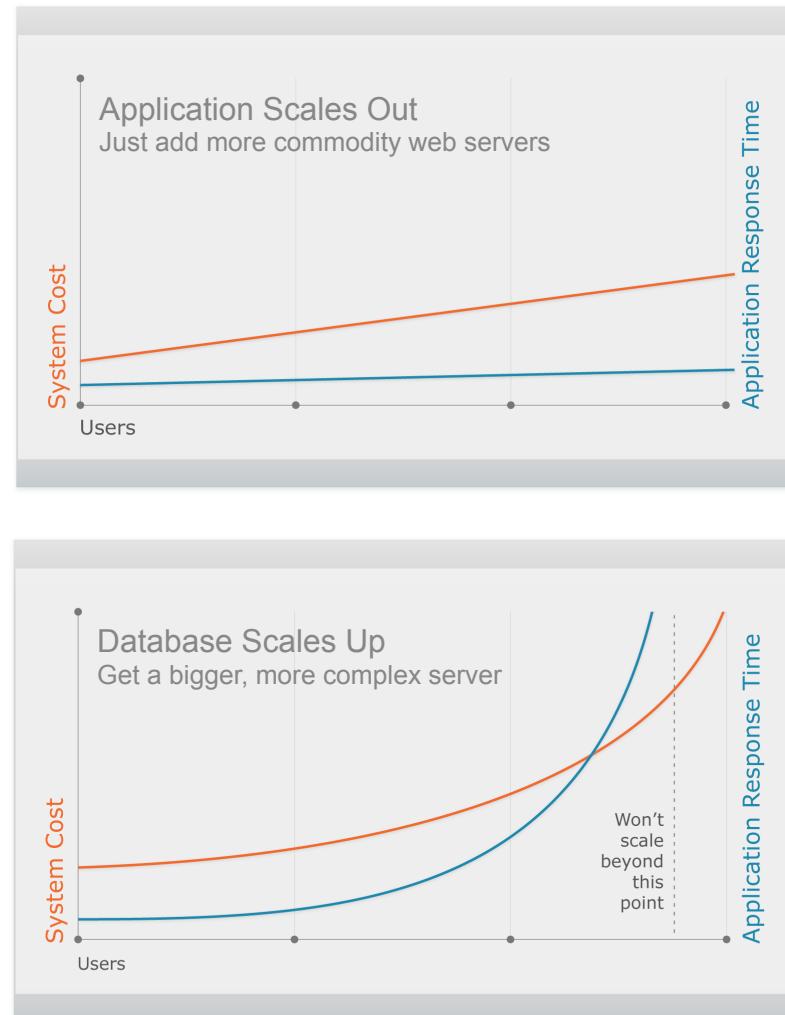
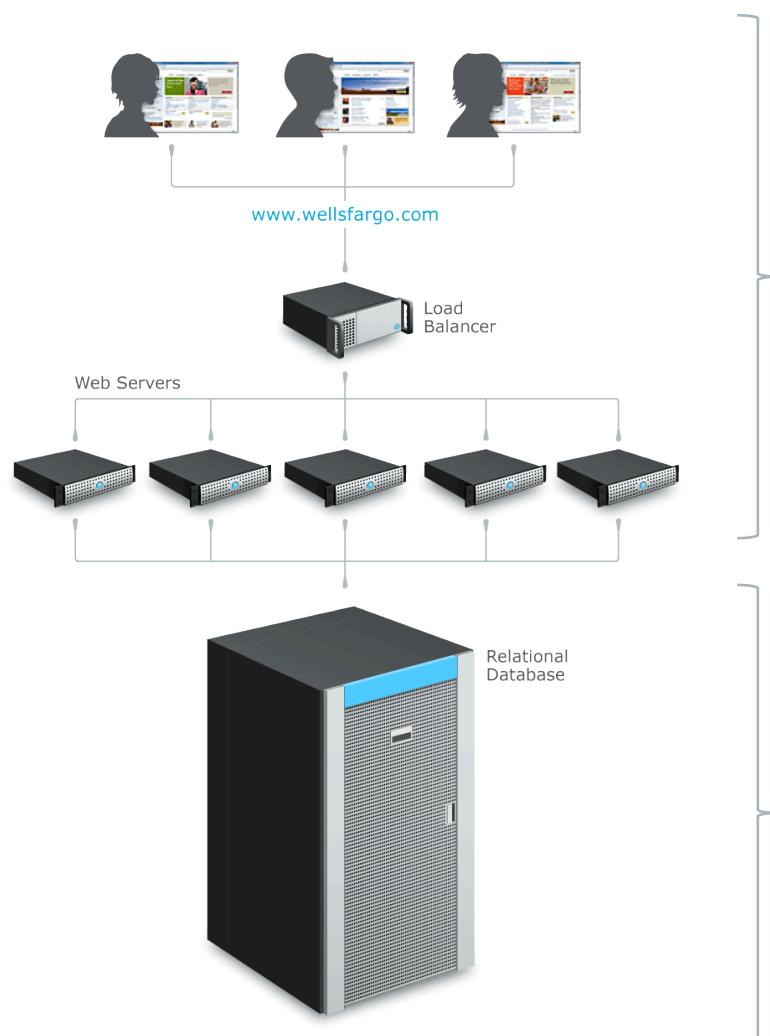


COUCHBASE

Changes in interactive software – NoSQL driver



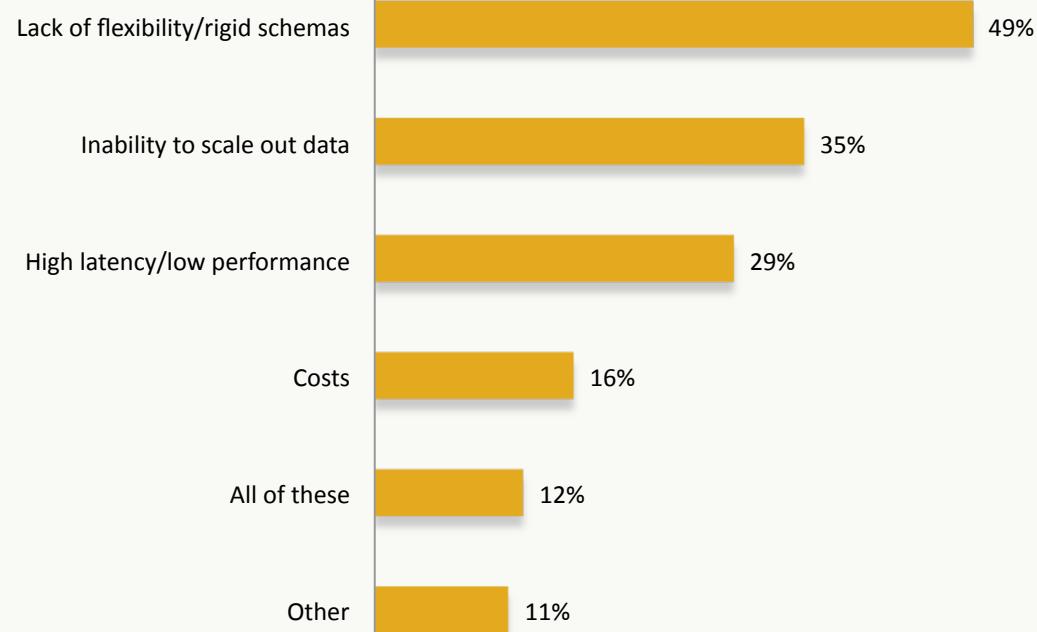
Modern interactive software architecture



Note – Relational database technology is great for what it is great for, but it is not great for this.

Survey: Schema inflexibility #1 adoption driver

What is the biggest data management problem driving your use of NoSQL in the coming year?



Source: Couchbase NoSQL Survey, December 2011, n=1351

Extending the scope of RDBMS technology

- Data partitioning (“sharding”)
 - Disruptive to reshuffle – impacts application
 - No cross-shard joins
 - Schema management on every shard
- Denormalizing
 - Increases speed
 - At the limit, provides complete flexibility
 - Eliminates relational query benefits
- Distributed caching
 - Accelerate reads
 - Scale out
 - Another tier, no write acceleration, coherency management

Lacking market solutions, users forced to invent



Bigtable
November 2006



Dynamo
October 2007



Cassandra
August 2008



Voldemort
February 2009

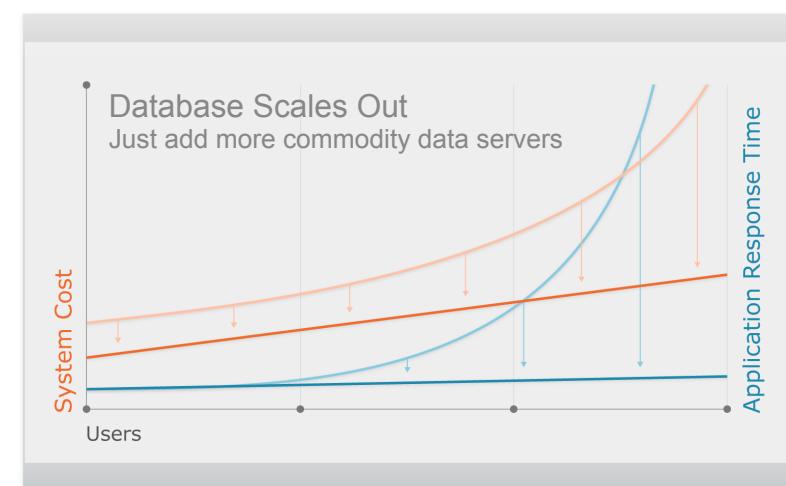
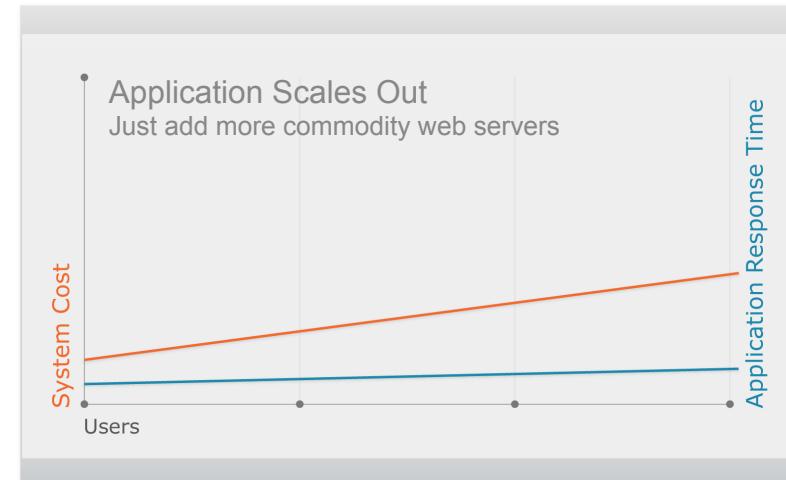
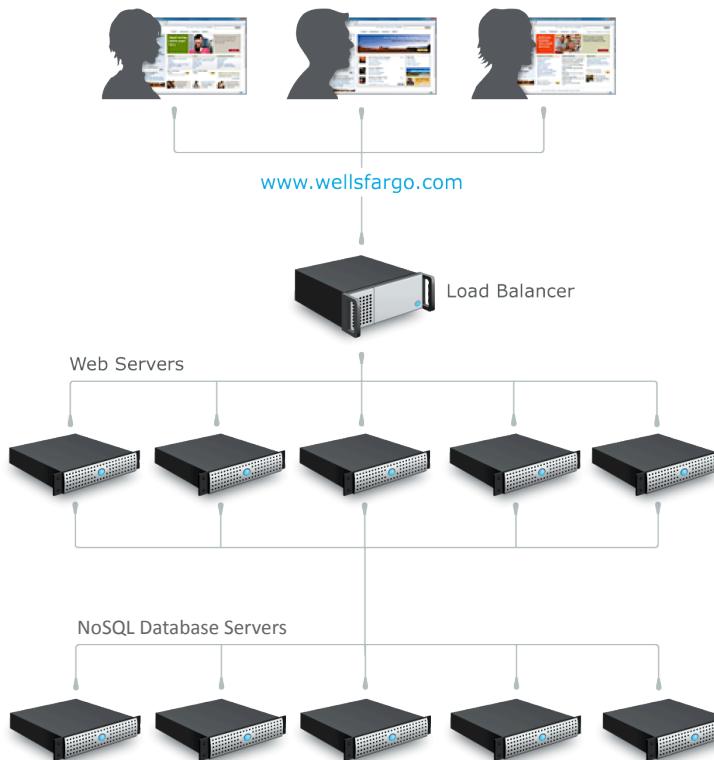
- No schema required before inserting data
- No schema change required to change data format
- Auto-sharding without application participation
- Distributed queries
- Integrated main memory caching
- Data synchronization (mobile, multi-datacenter)



Very few organizations want to (fewer can) build and maintain database software technology.
But every organization building interactive web applications needs this technology.

NoSQL database matches application logic tier architecture

Data layer now scales with linear cost and constant performance.



▶ Scaling out flattens the cost *and* performance curves.

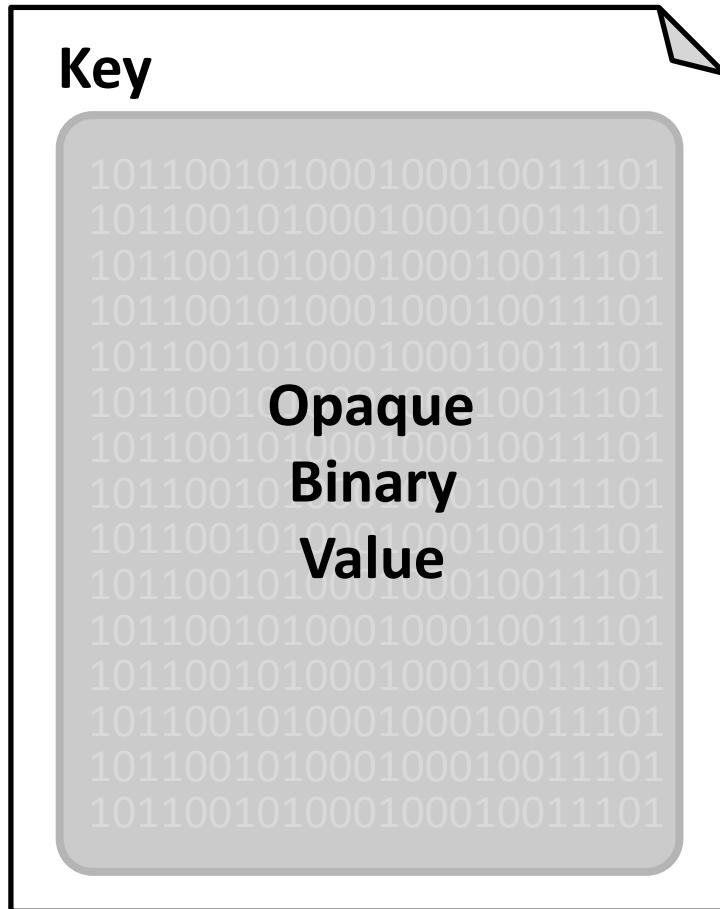
NOSQL TAXONOMY

The Key-Value Store – the foundation of NoSQL

Key

101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
**Opaque
Binary
Value**
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101
101100101000100010011101

Memcached – the NoSQL precursor



memcached

In-memory only

Limited set of operations

Blob Storage: Set, Add, Replace, CAS

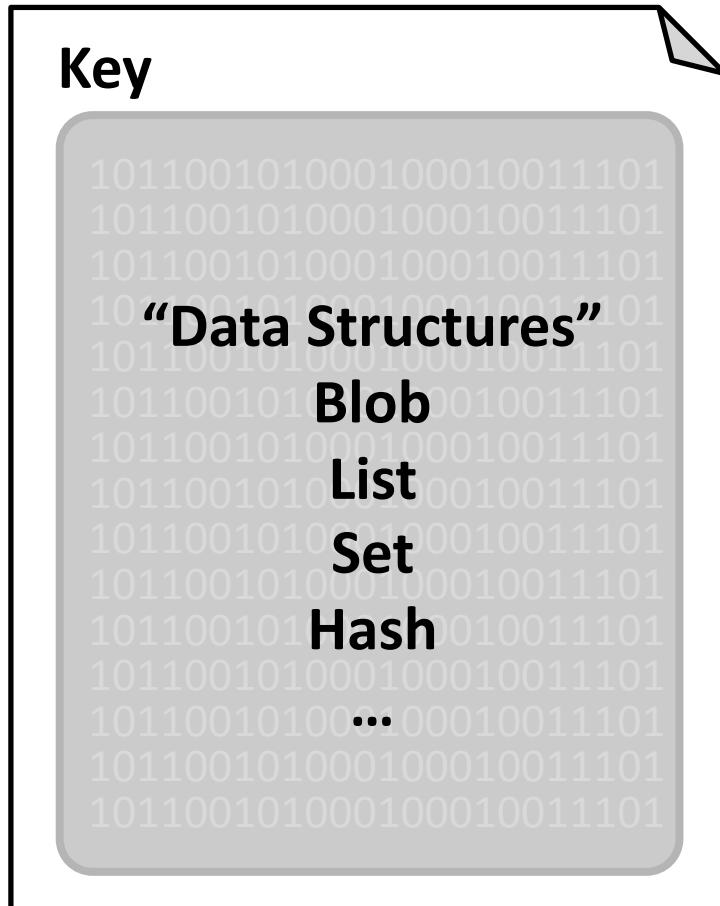
Retrieval: Get

Structured Data: Append, Increment

“Simple and fast.”

Challenges: cold cache, disruptive elasticity

Redis – More “Structured Data” commands



redis

In-memory only

Vast set of operations

Blob Storage: Set, Add, Replace, CAS

Retrieval: Get, Pub-Sub

Structured Data: Strings, Hashes, Lists, Sets,
Sorted lists

Example operations for a Set

Add, count, subtract sets, intersection, is member?, atomic move from one set to another

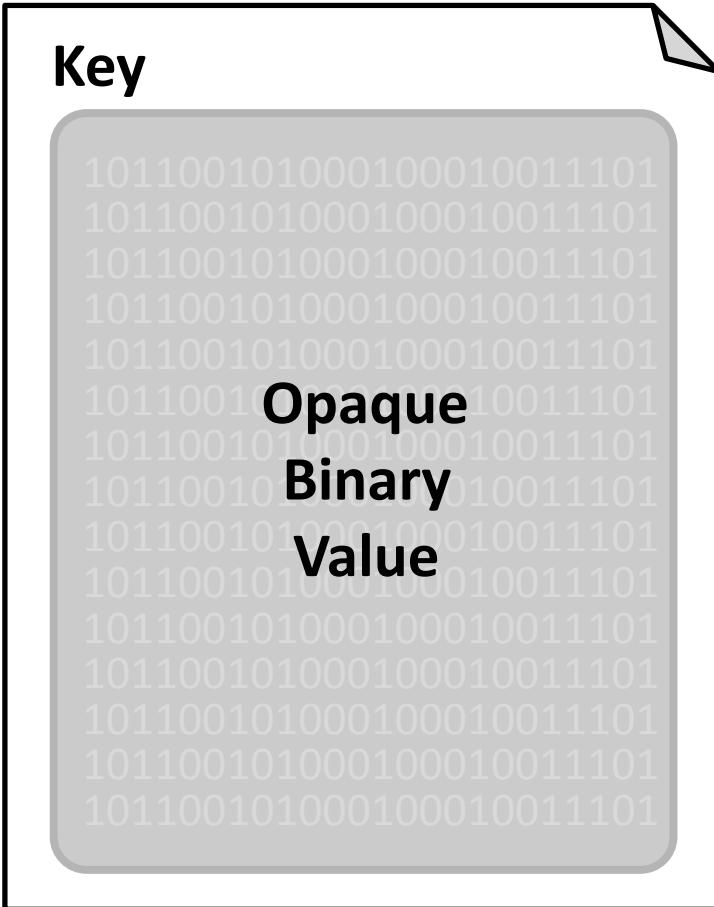
NoSQL catalog



Key-Value	Data Structure	Document	Column	Graph
 memcached	 redis			

Cache
(memory only)

Membase – From key-value cache to database



membase

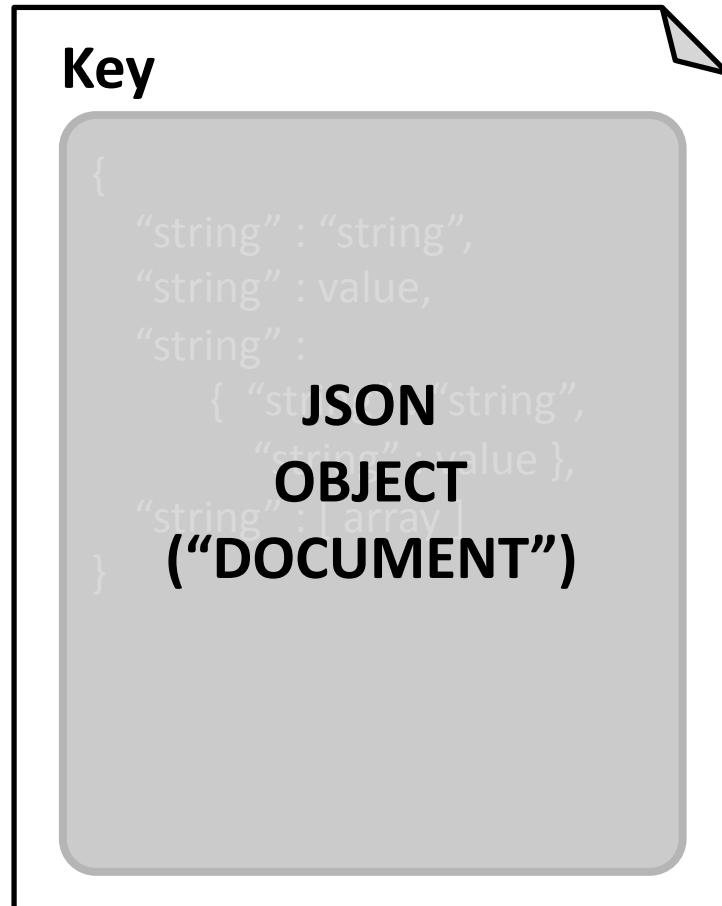
Disk-based with built-in memcached cache
Cache refill on restart
Memcached compatible (drop in replacement)
Highly-available (data replication)
Add or remove capacity to live cluster

“Simple, fast, elastic.”

NoSQL catalog

	Key-Value	Data Structure	Document	Column	Graph
Cache (memory only)	 memcached	 redis			
Database (memory/disk)					 membase

Couchbase – document-oriented database



Couchbase

Auto-sharding

Disk-based with built-in memcached cache

Cache refill on restart

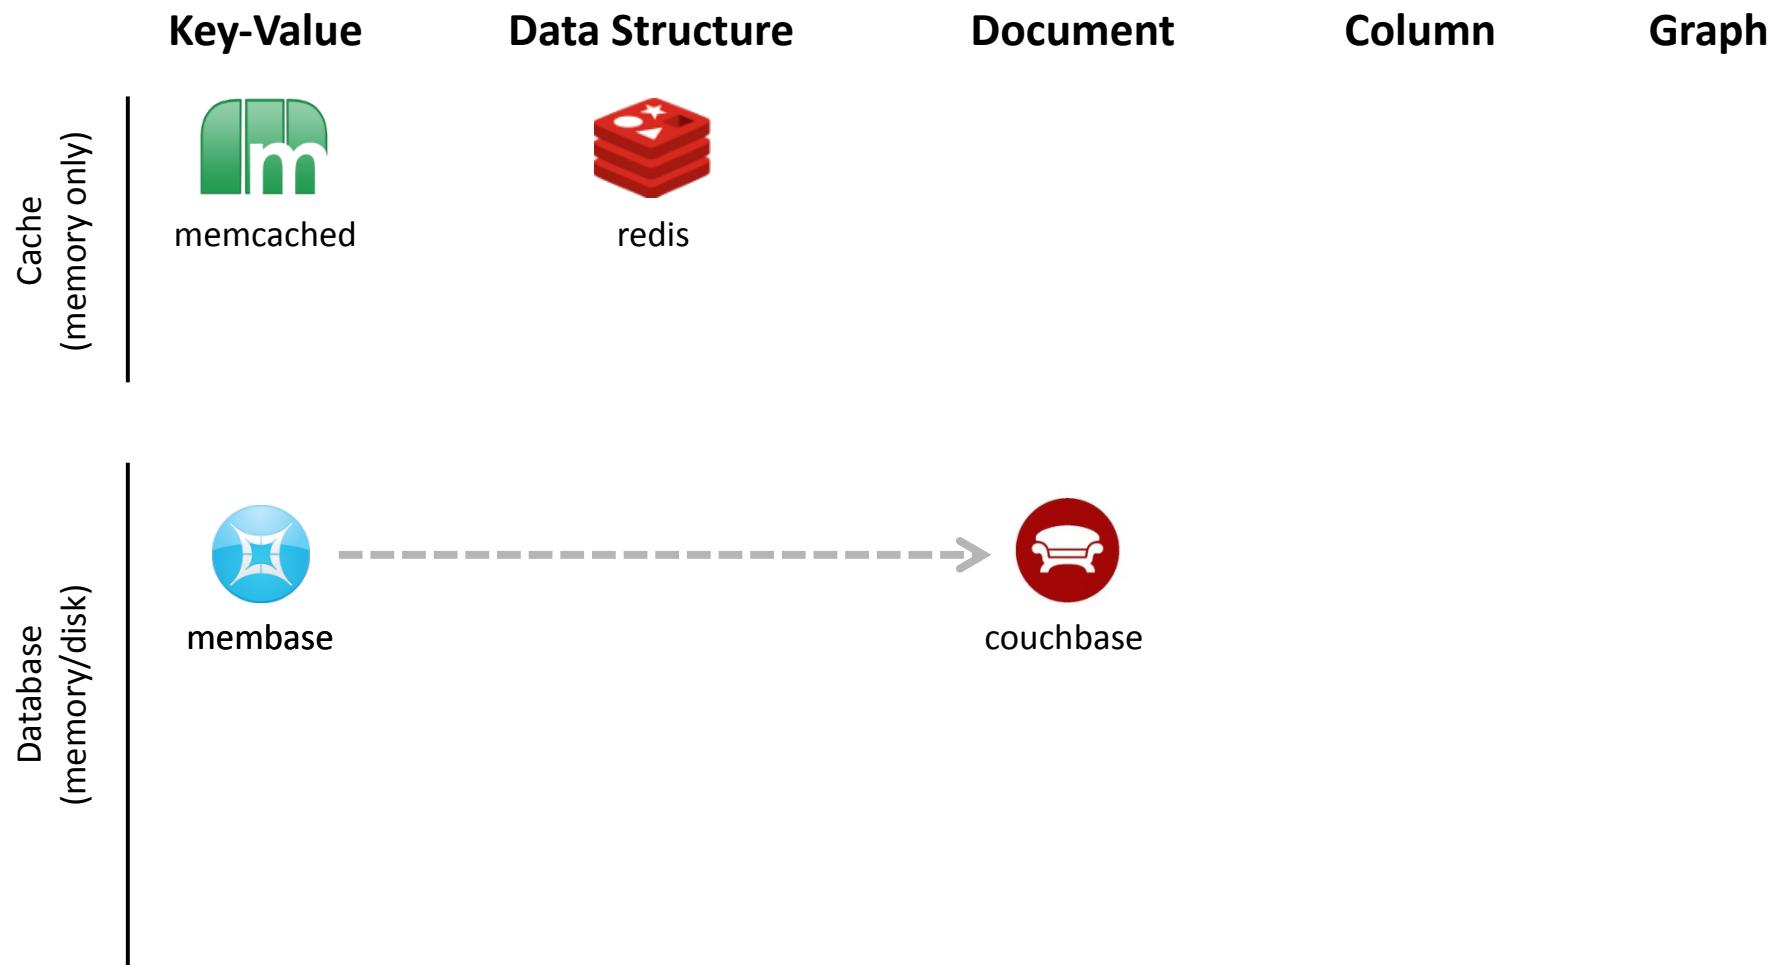
Memcached compatible (drop in replace)

Highly-available (data replication)

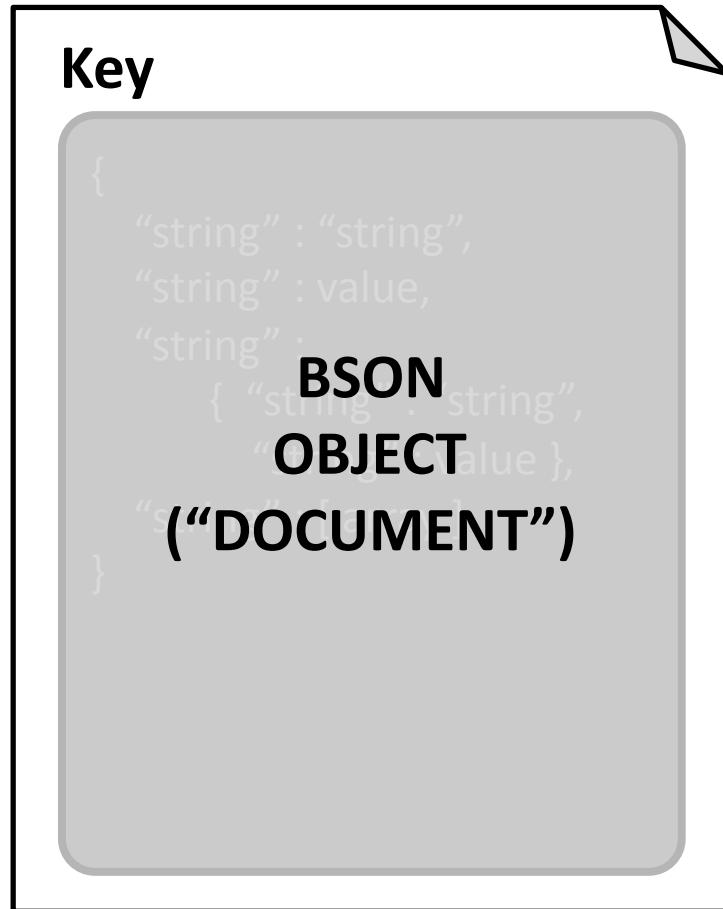
Add or remove capacity to live cluster

When values are JSON objects (“documents”):
Create indices, views and query against the
views

NoSQL catalog



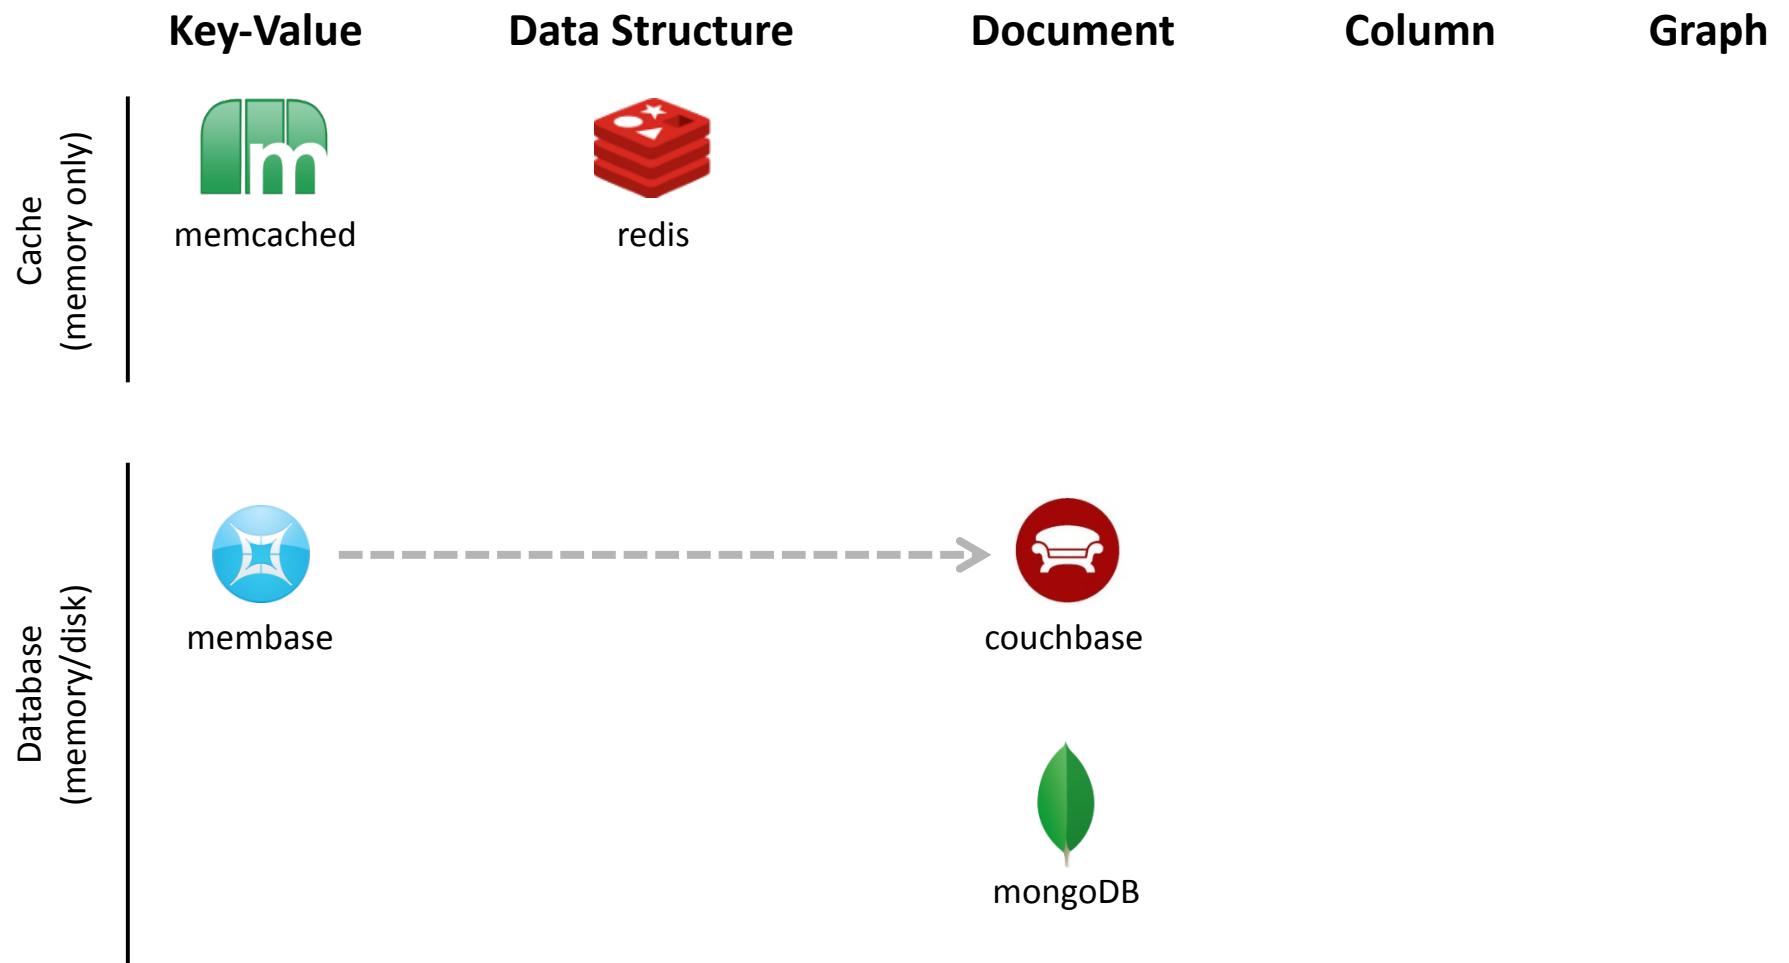
MongoDB – Document-oriented database



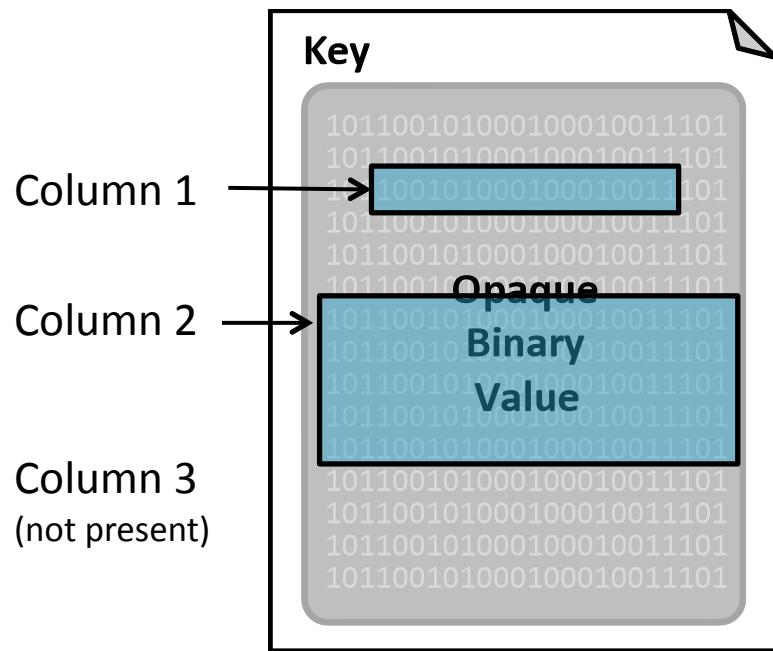
MongoDB

Disk-based with in-memory “caching”
BSON (“binary JSON”) format and wire protocol
Master-slave replication
Auto-sharding
Values are BSON objects
Supports ad hoc queries – best when indexed

NoSQL catalog



Cassandra – Column overlays



Cassandra

Disk-based system

Clustered

External caching required for low-latency reads

“Columns” are overlaid on the data

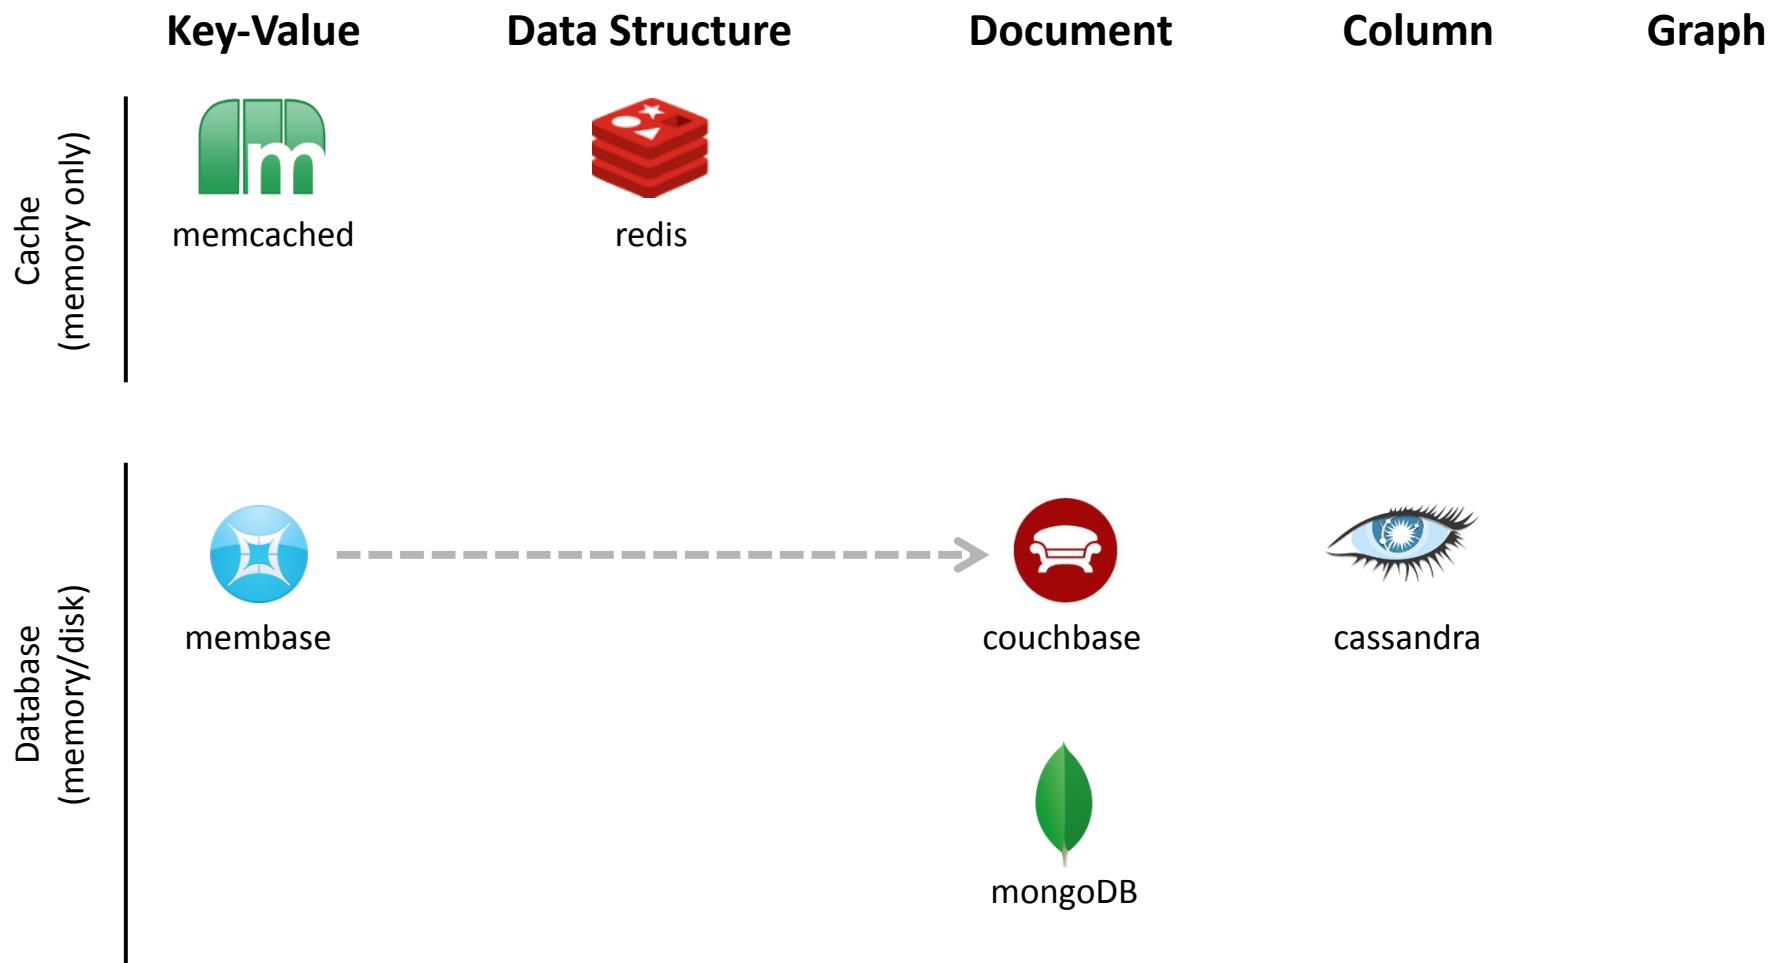
Not all rows must have all columns

Supports efficient queries on columns

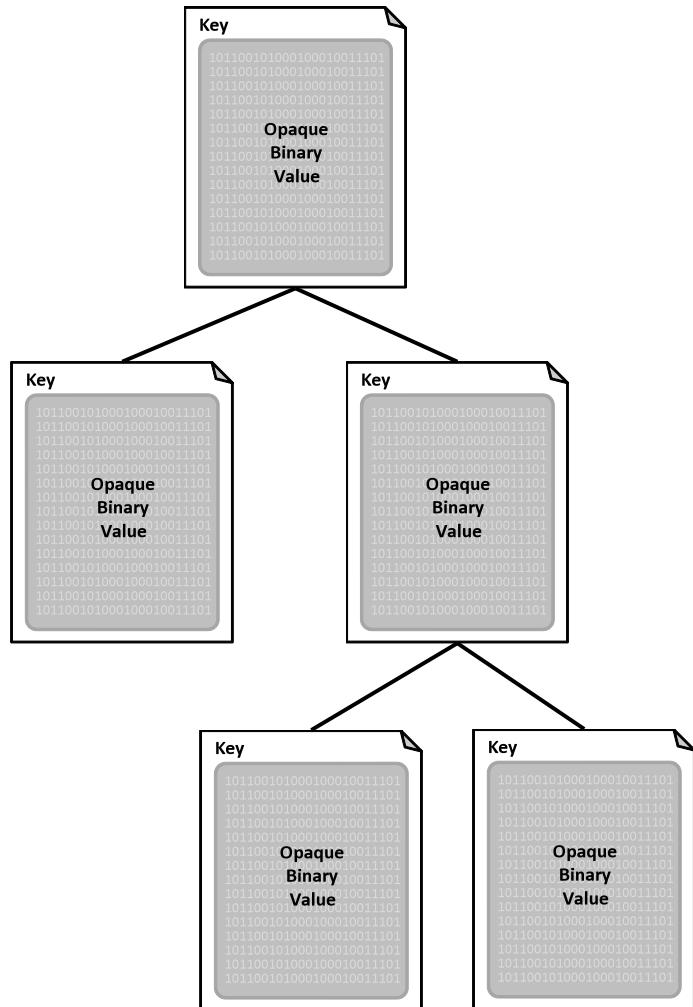
Restart required when adding columns

Good cross-datacenter support

NoSQL catalog



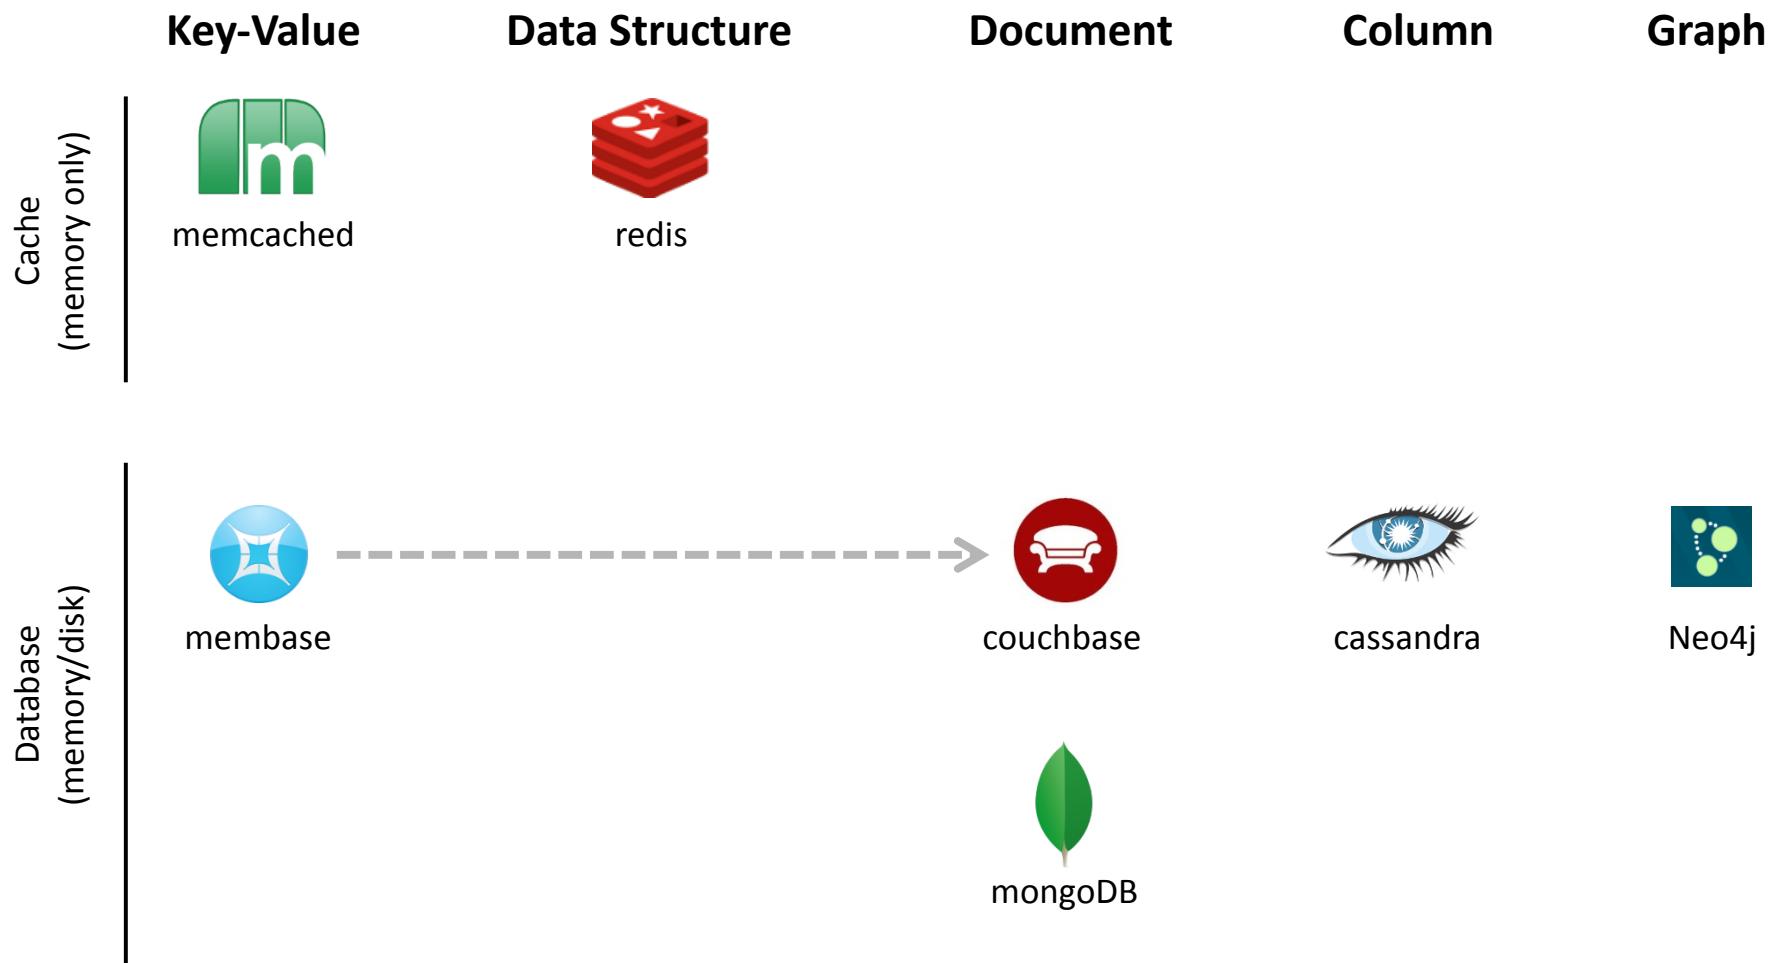
Neo4j – Graph database



Neo4j

Disk-based system
External caching required for low-latency reads
Nodes, relationships and paths
Properties on nodes
Delete, Insert, Traverse, etc.

NoSQL catalog



Document-oriented database advantages

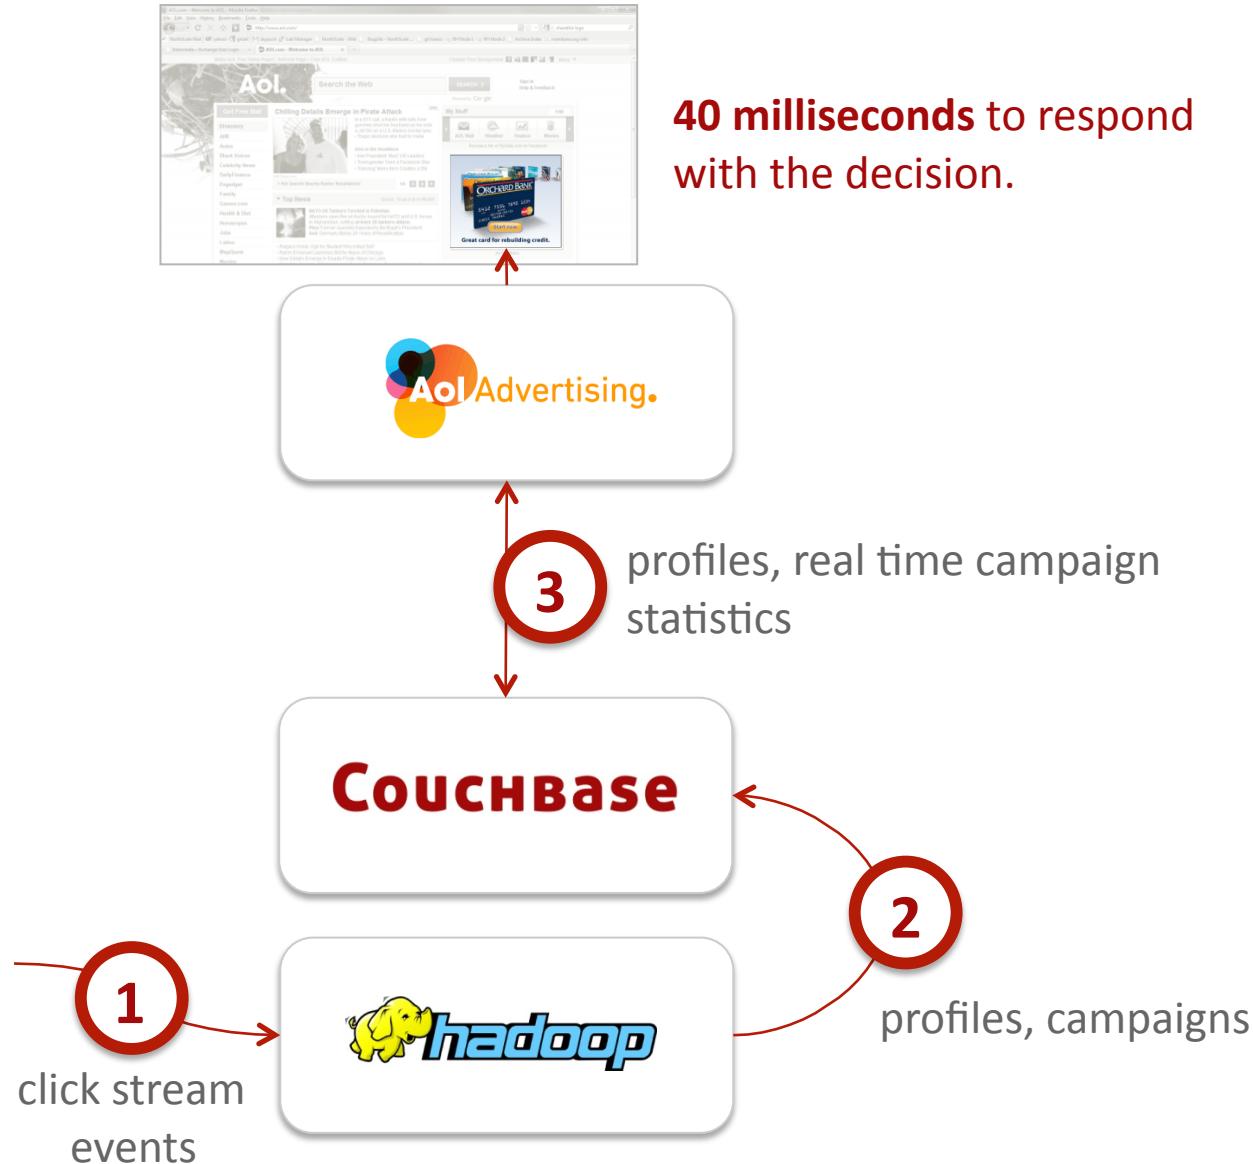
Performance. The document data model keeps related data in a single physical location in memory and on disk (a document). This allows consistently low-latency access to the data – reads and writes happen with very little delay. Database latency can result in perceived “lag” by the player of a game and avoiding it is a key success criterion.

Dynamic elasticity. Because the document approach keeps records “in one place” (a single document in a contiguous physical location), it is much easier to move the data from one server to another while maintaining consistency – and without requiring any game downtime. Moving data between servers is required to add and remove cluster capacity to cost-effectively match the aggregate performance needs of the application to the performance capability of the database. Doing this at any time without stopping the revenue flow of the game can make a material difference in game profitability.

Schema flexibility. While all NoSQL databases provide schema flexibility. Key-value and document-oriented databases enjoy the most flexibility. Column-oriented databases still require maintenance to add new columns and to group them. A key-value or document-oriented database requires no database maintenance to change the database schema (to add and remove “fields” or data elements from a given record).

Query flexibility. Balancing schema flexibility with query expressiveness (the ability to ask the database questions, for example “return me a list of all the farms in which a player purchased a black sheep last month”) is important. While a key-value database is completely flexible, allowing a user to put any desired value in the “value” part of the key-value pair, it doesn’t provide the ability to ask questions. It only permits accessing the data record associated with a given key. I can ask for the farm data for user A, B, C ... and see if they have a black sheep, but I can’t ask the database to do that work on my behalf. Document-databases provide the best balance of schema flexibility without giving up the ability to do sophisticated queries.

Big data and NoSQL – Hadoop and Couchbase



COUCHBASE

Couchbase Server



Simple. Fast. Elastic. NoSQL.

Couchbase automatically distributes data across commodity servers. Built-in caching enables apps to read and write data with sub-millisecond latency. And with no schema to manage, Couchbase effortlessly accommodates changing data management requirements.



Cache



24x7x365



Clone to Grow



Production Ready



Reliable



Replication

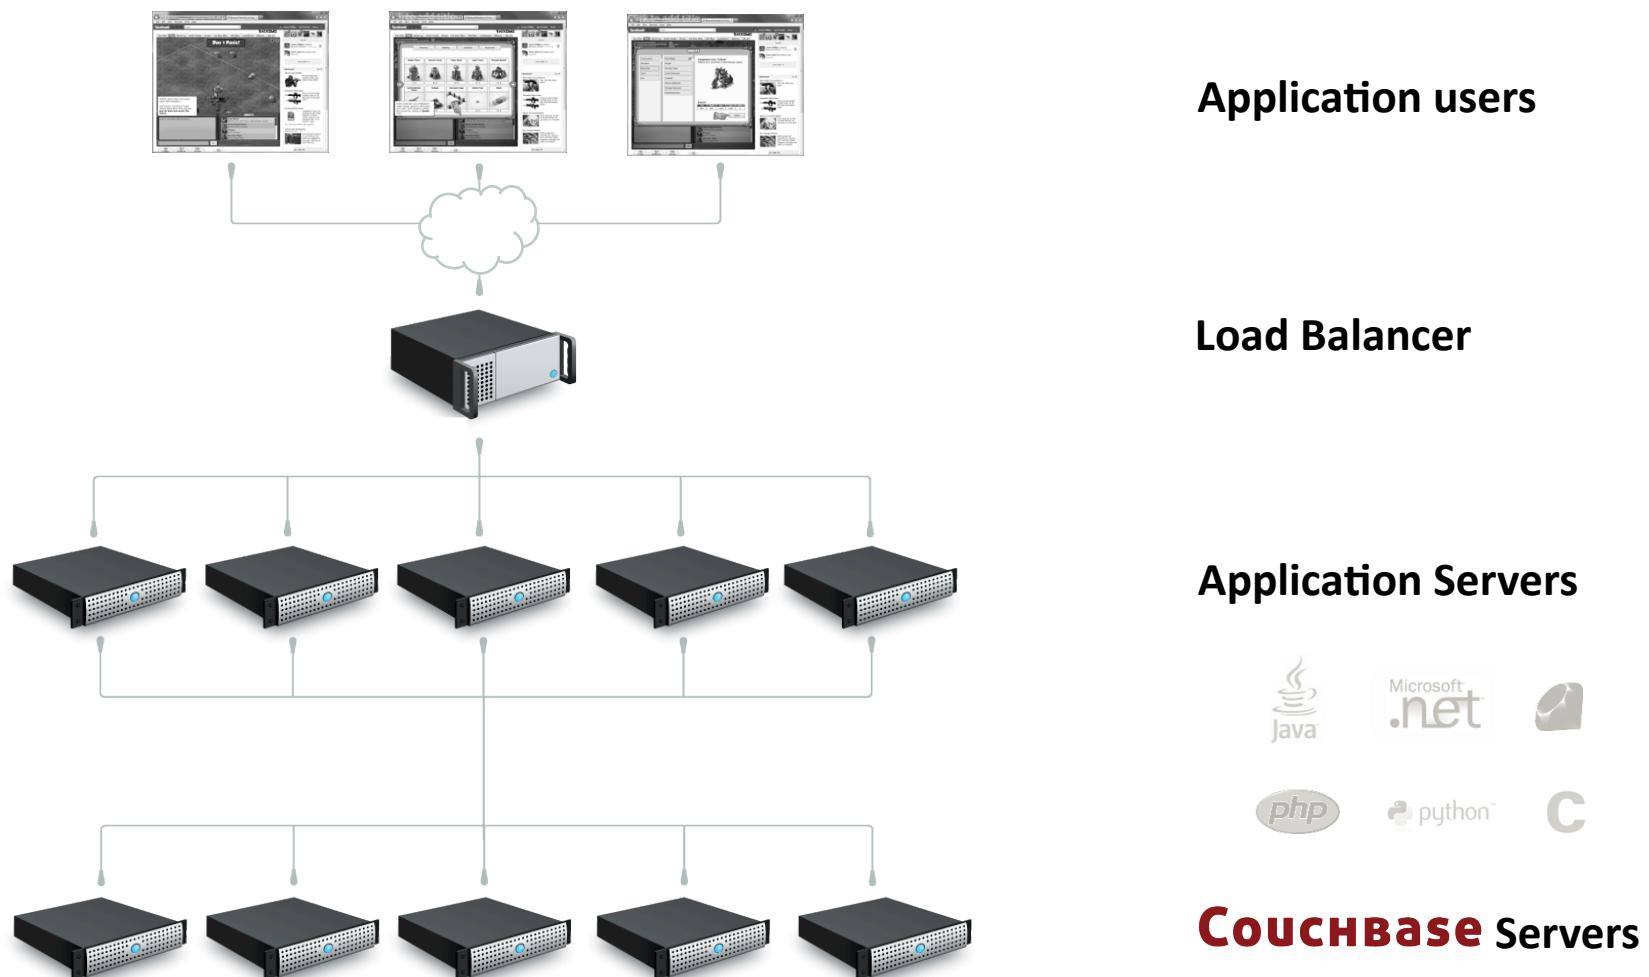


SDKs

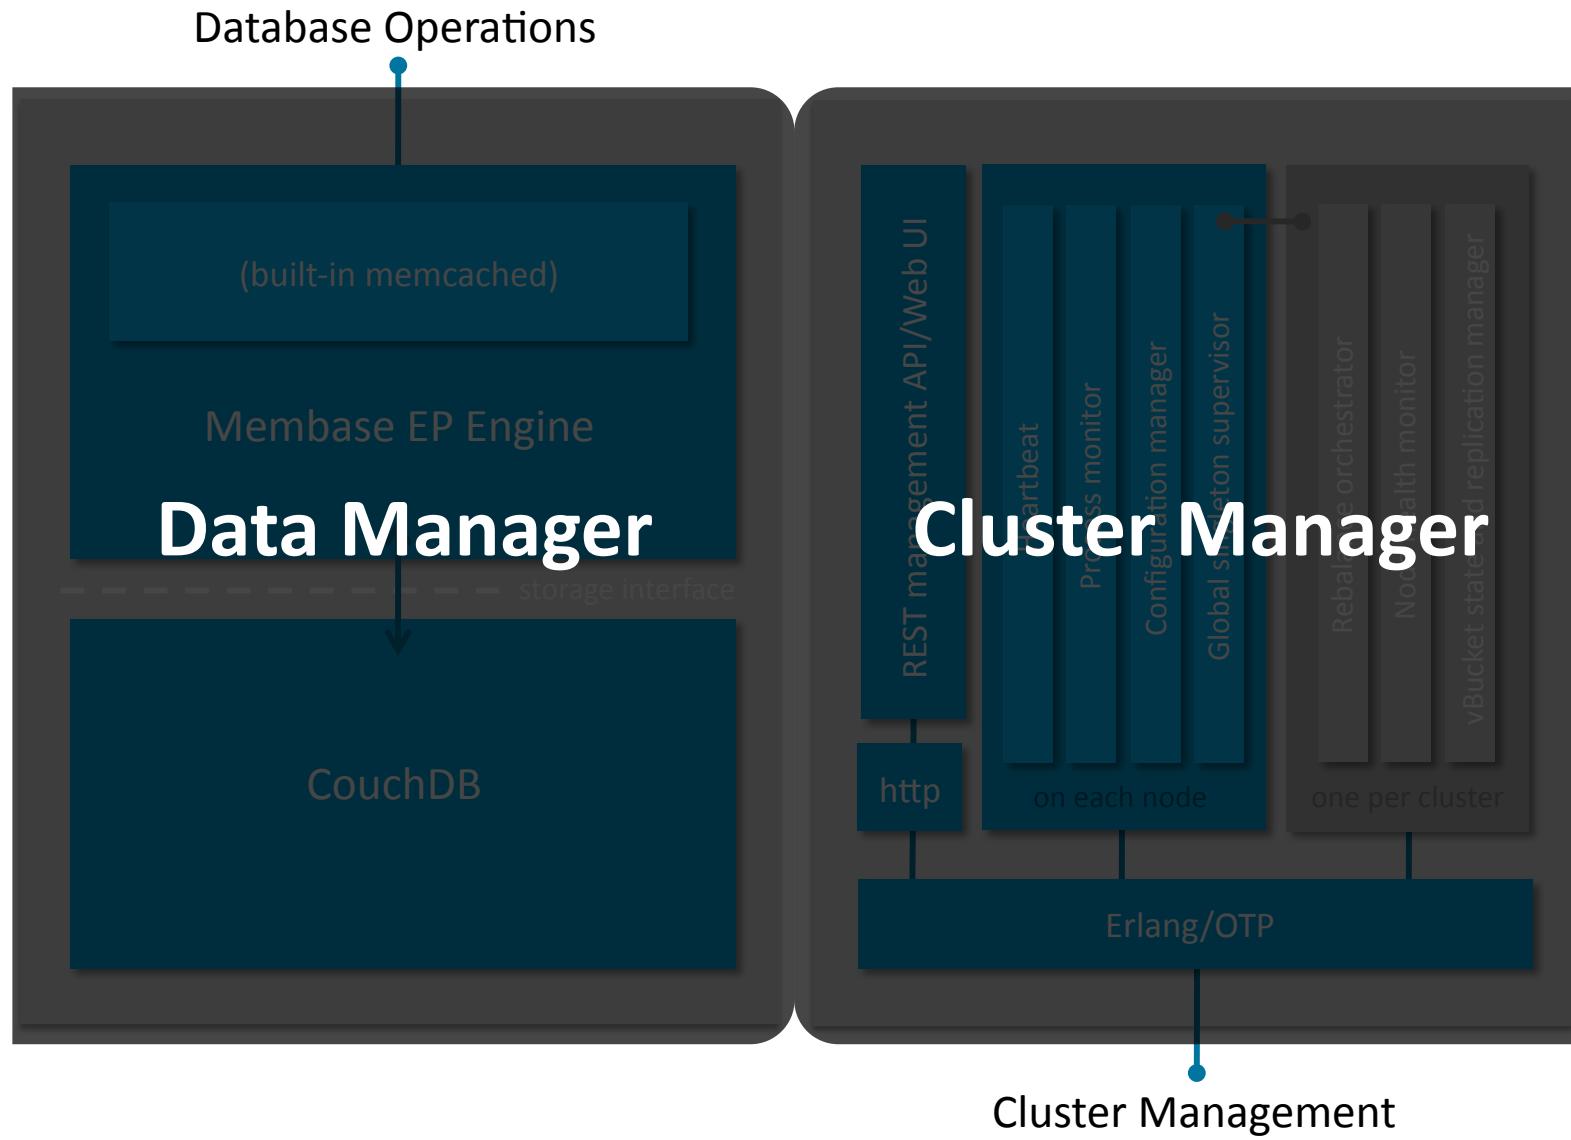
Representative user list



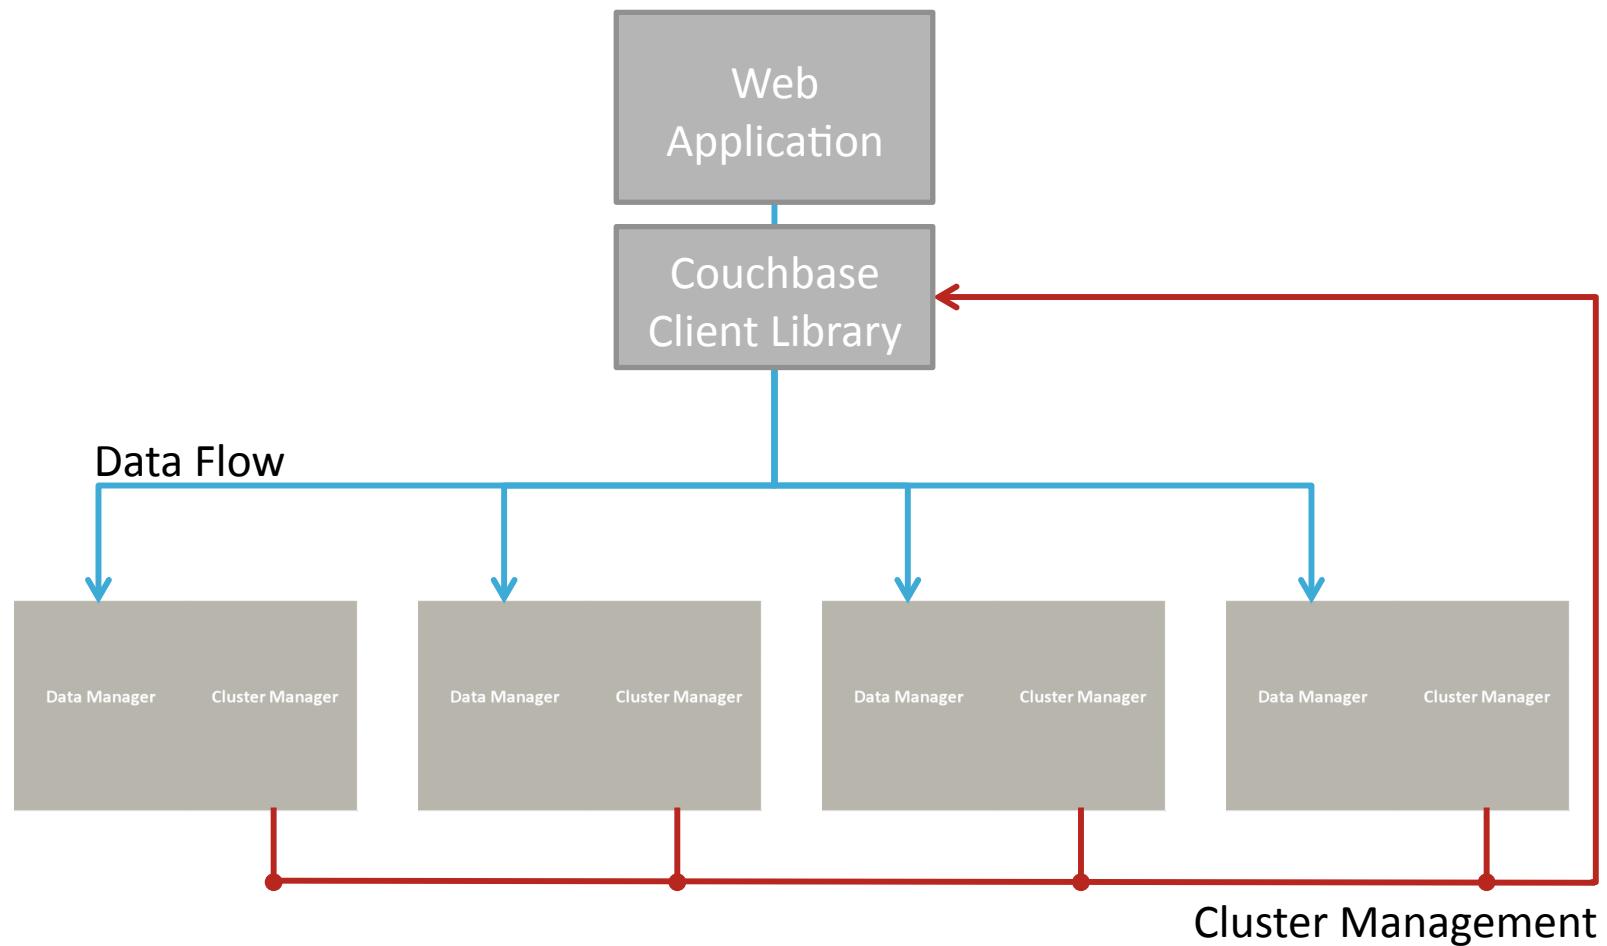
Typical Couchbase production environment



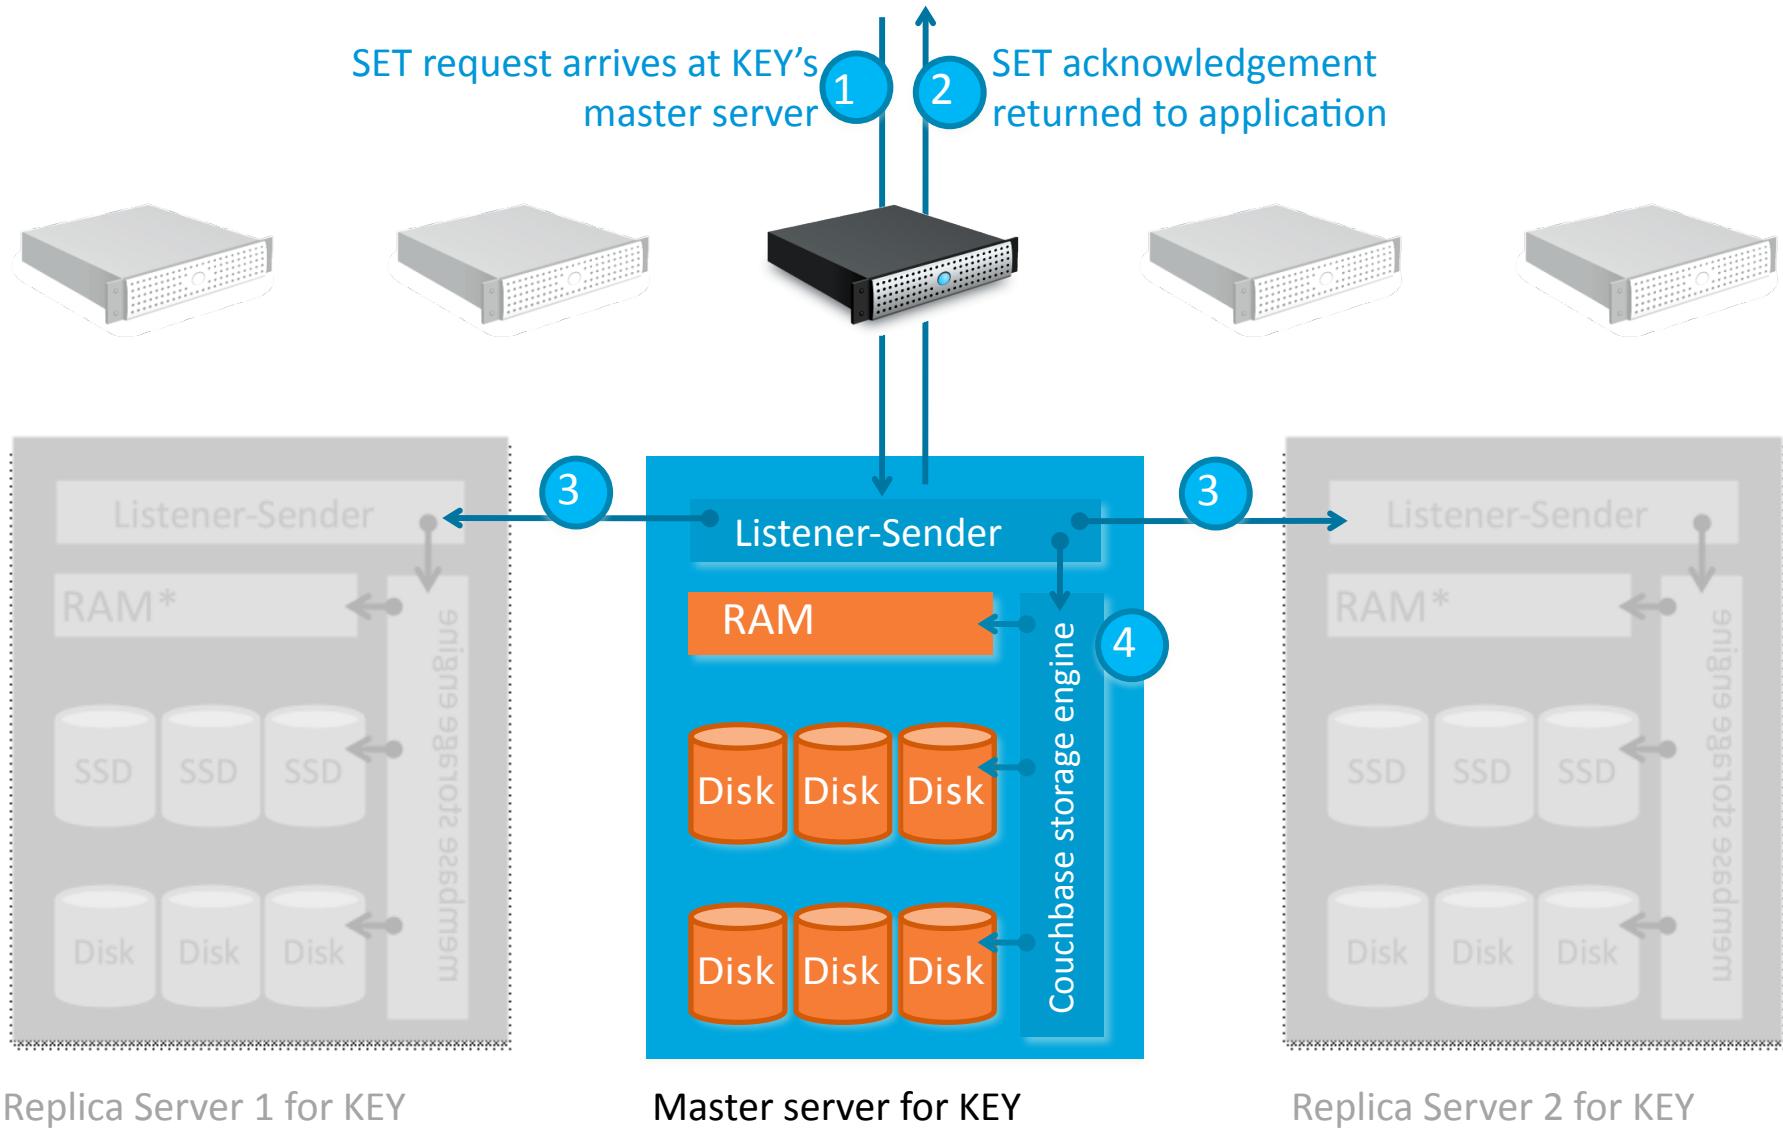
Couchbase architecture



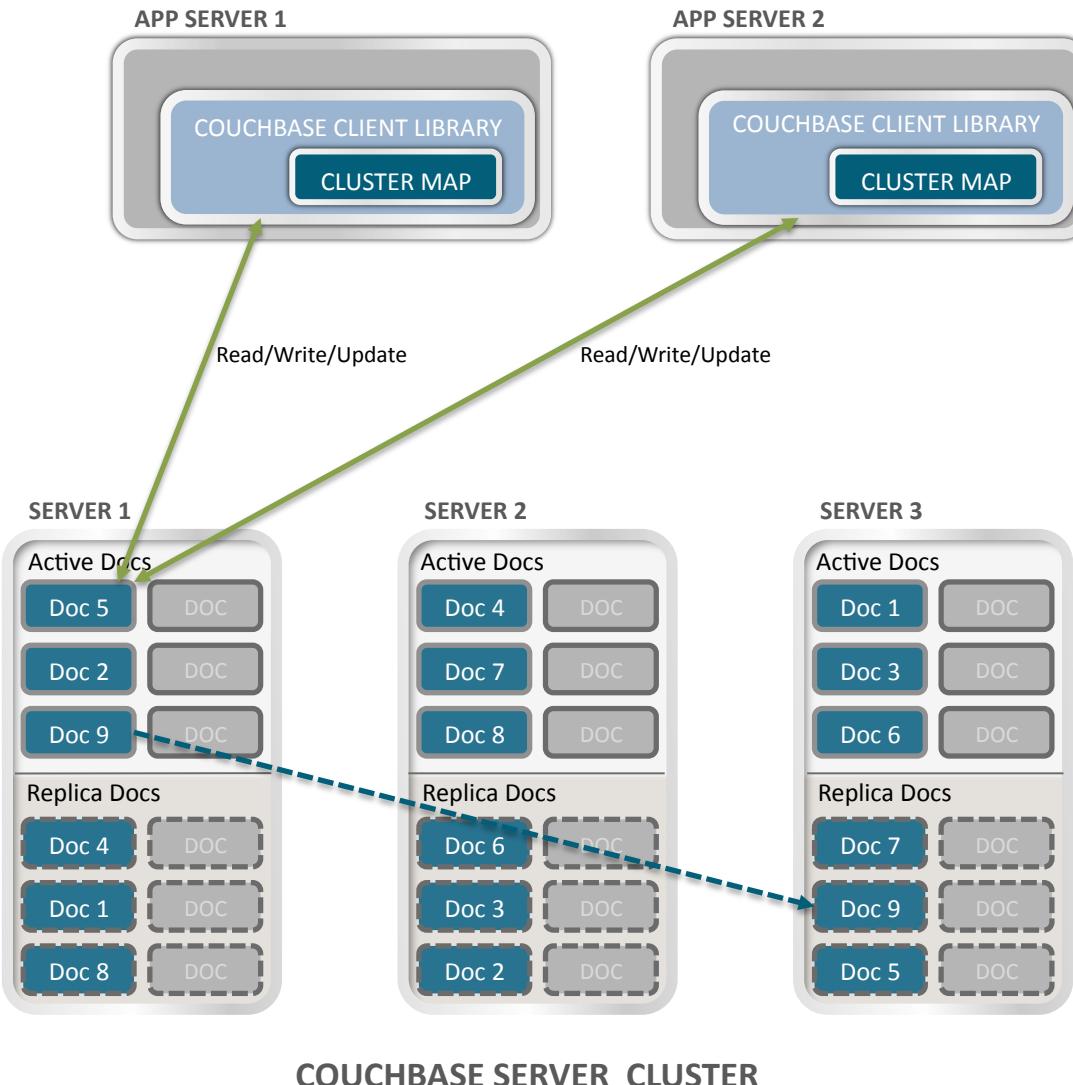
Couchbase deployment



Clustering With Couchbase

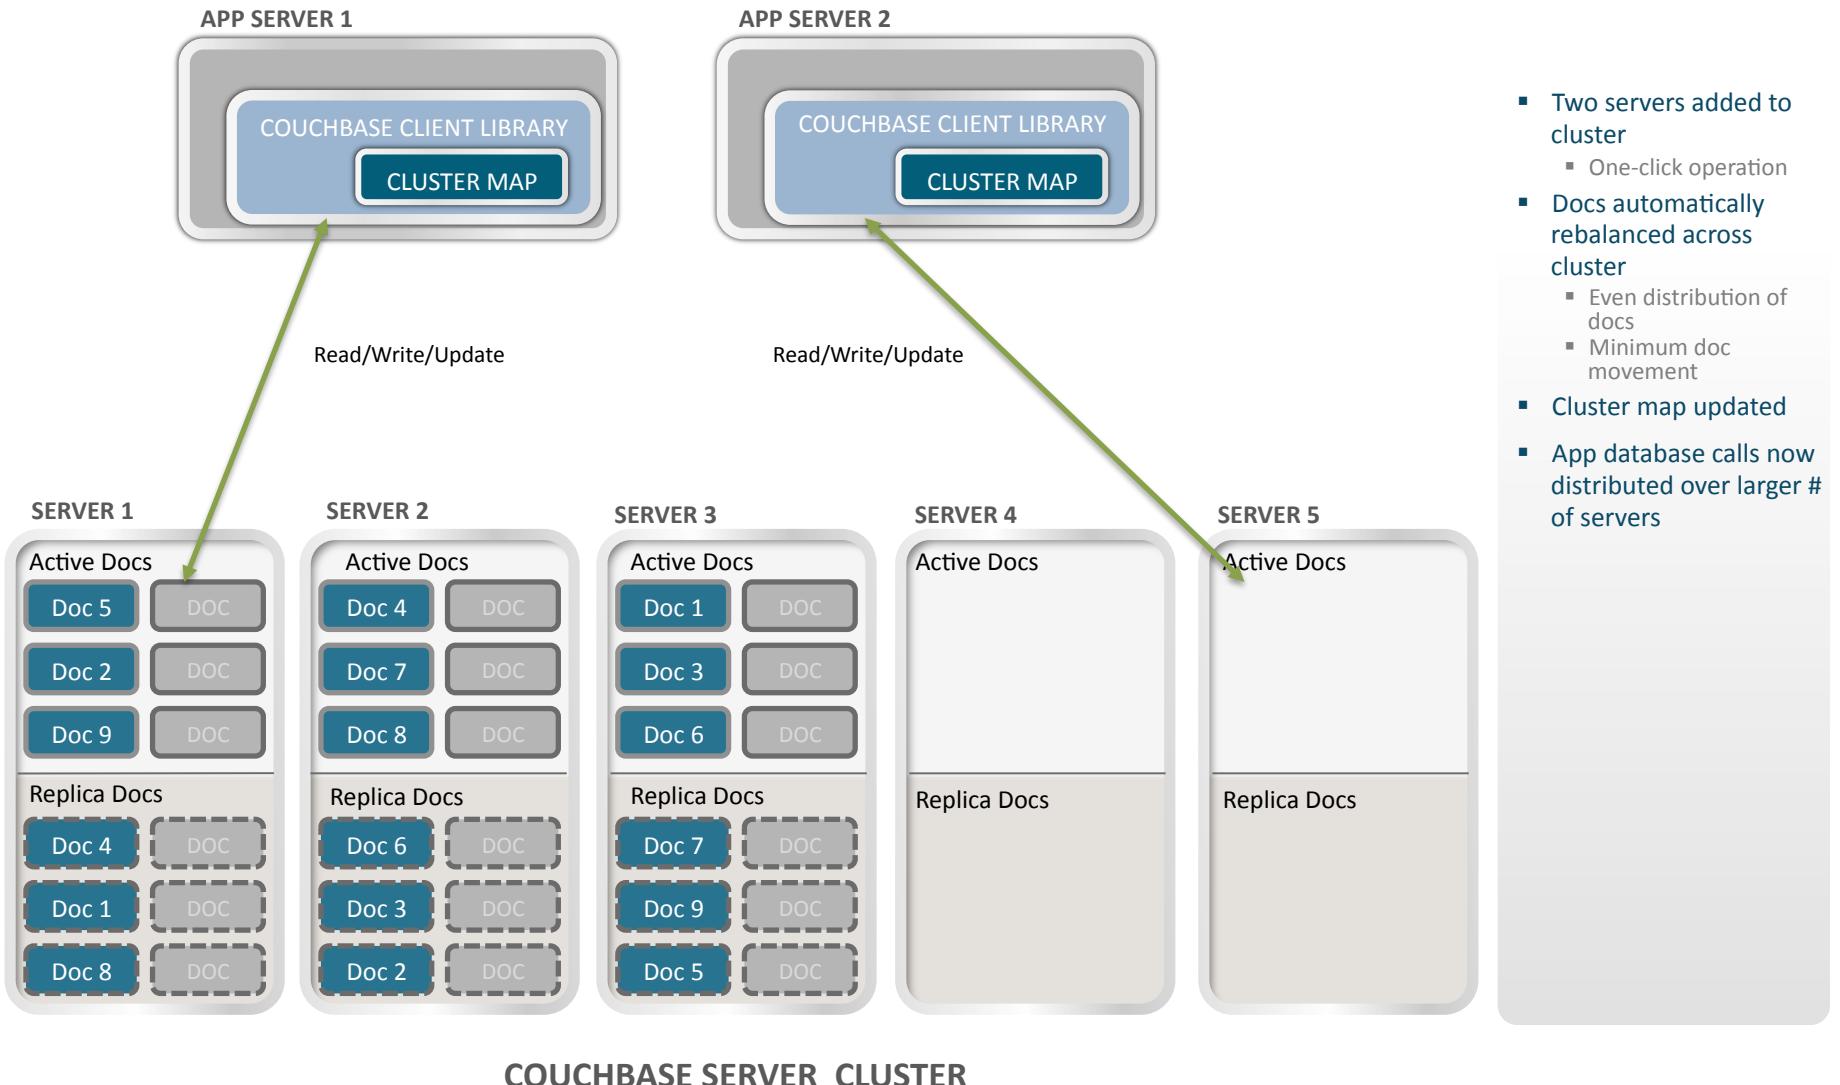


Basic Operation



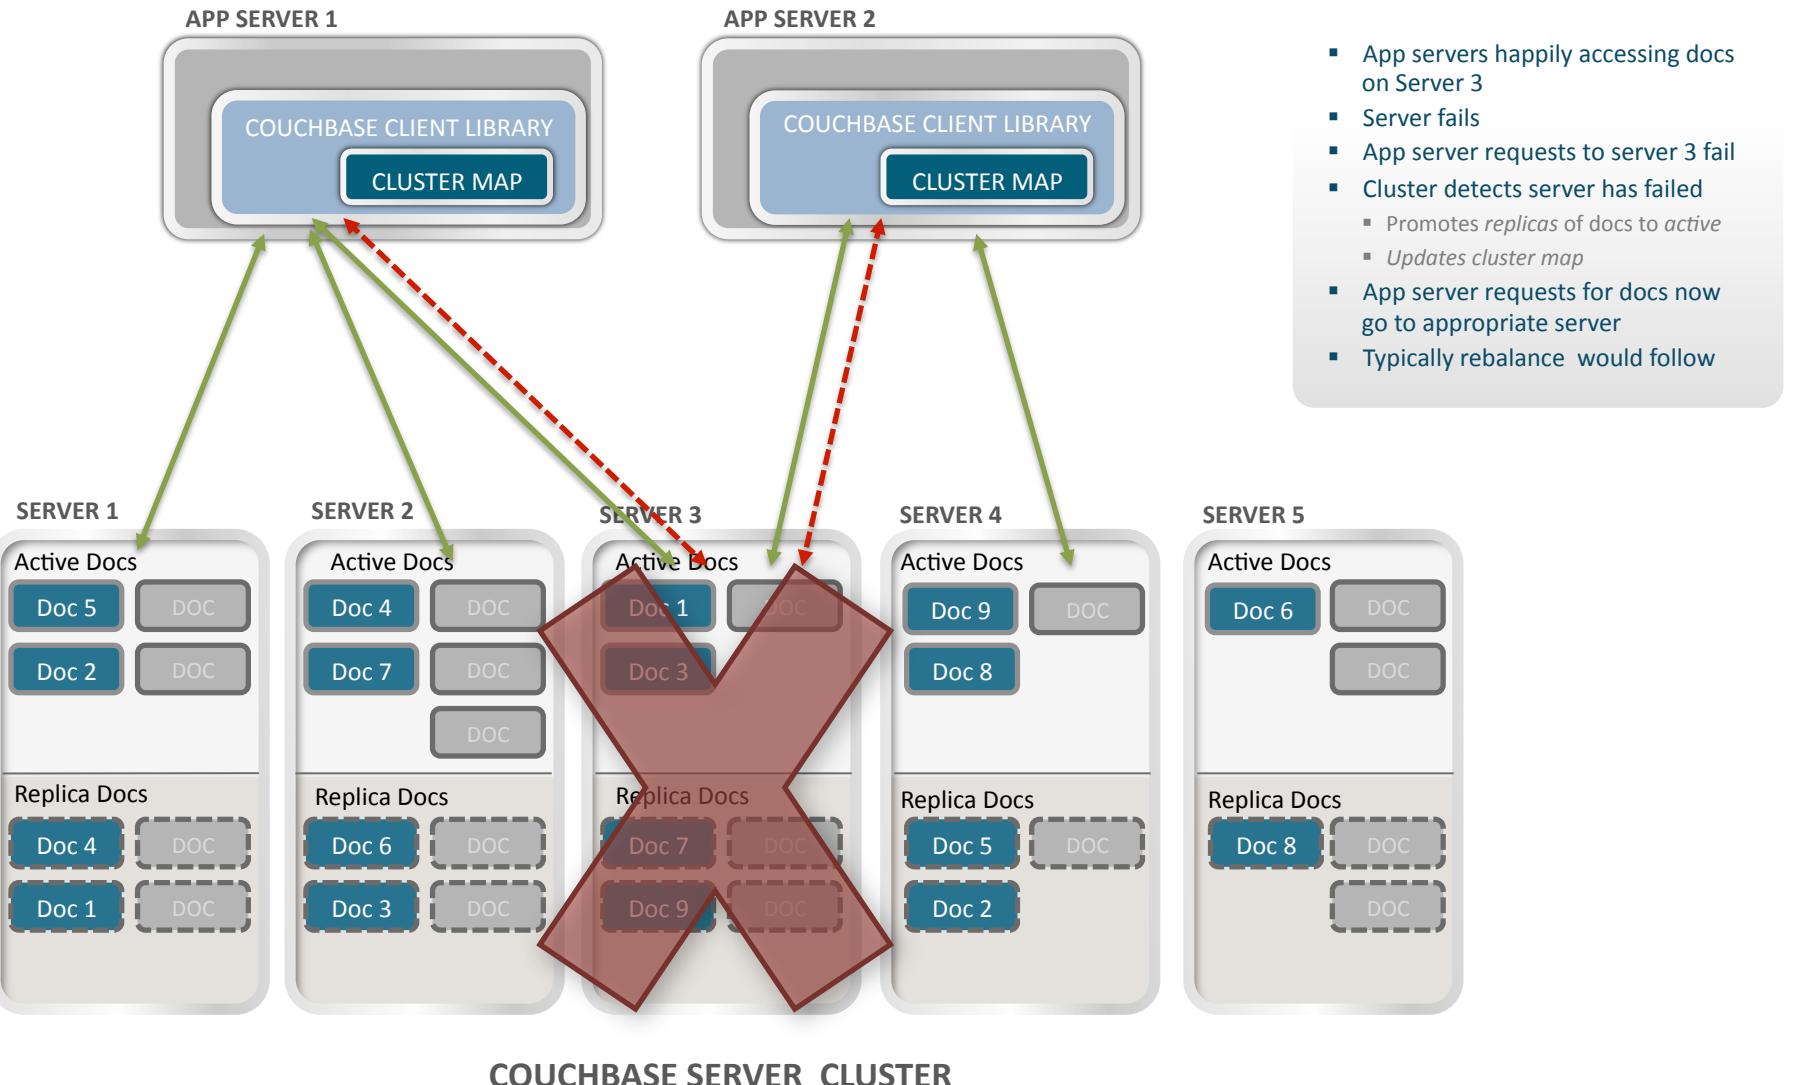
- Docs distributed evenly across servers in the cluster
- Each server stores both *active* & *replica* docs
 - Only one server active at a time
- Client library provides app with simple interface to database
- Cluster map provides map to which server doc is on
 - App never needs to know
- App reads, writes, updates docs
- Multiple App Servers can access same document at same time

Add Nodes



User Configured Replica Count = 1

Fail Over Node



User Configured Replica Count = 1

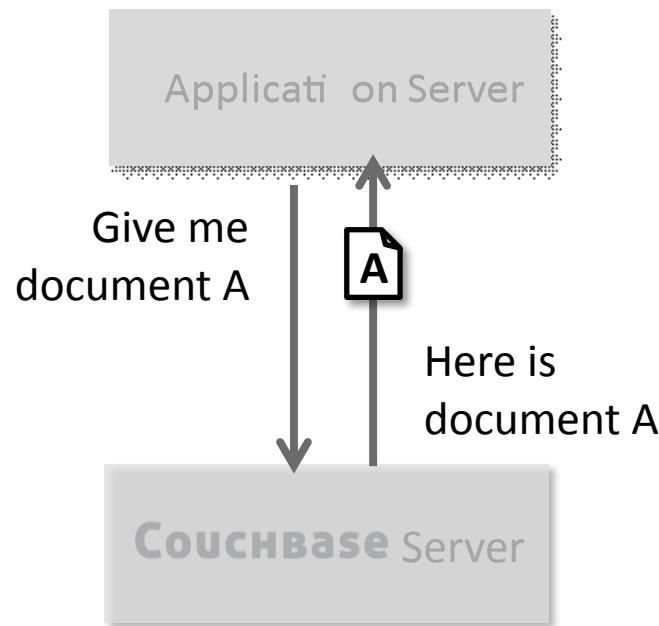


COUCHBASE SOLUTION OPERATING A CLUSTER

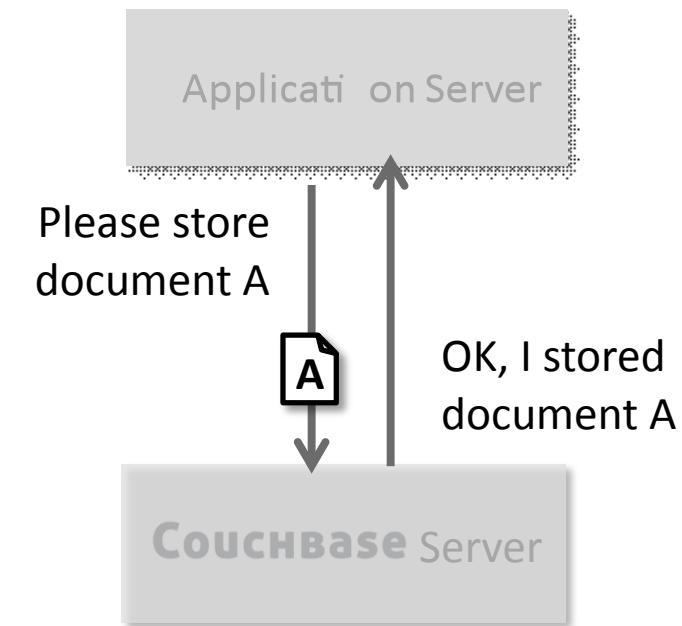
Reading and Writing



Reading Data

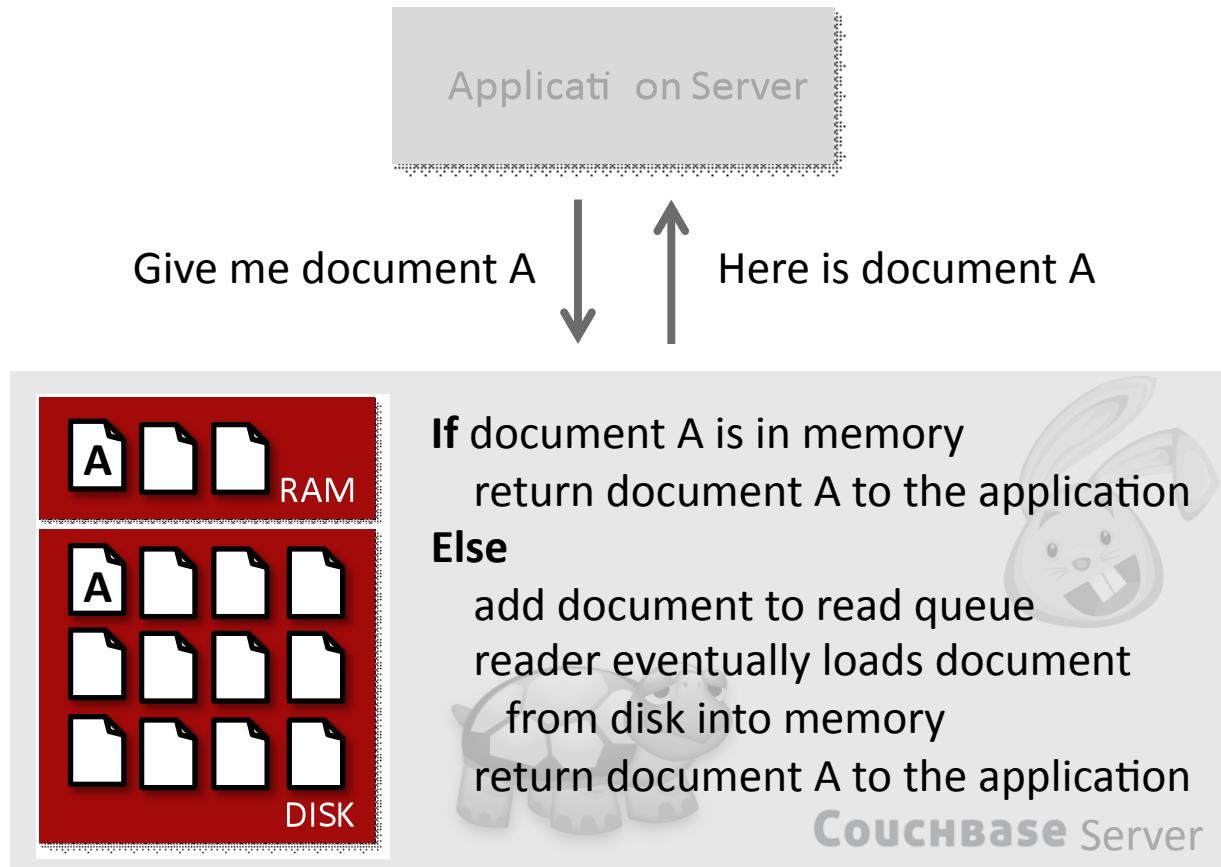


Writing Data

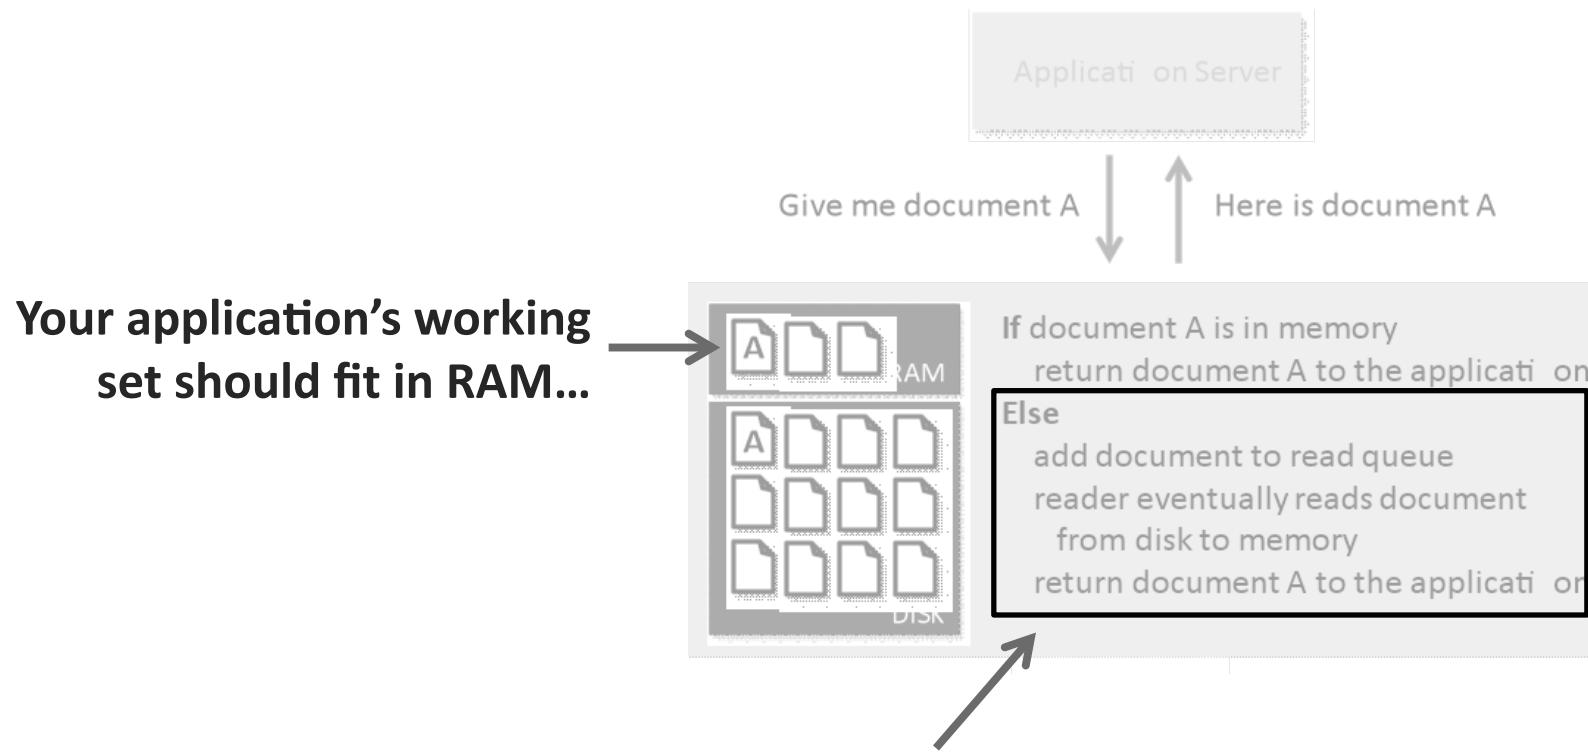


(We'll save the arithmetic for the sizing section :)

Reading data



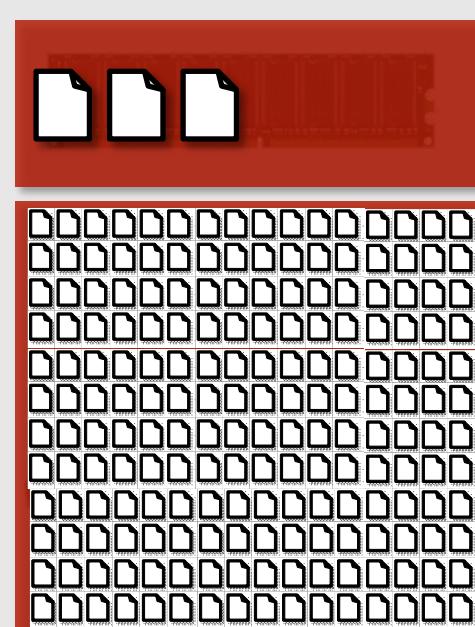
Keeping working data set in RAM is key to read performance



... or else! (because you don't want the “else” part happening very often – it is MUCH slower than a memory read and you could have to wait in line an indeterminate amount of time for the read to happen.)

Working set ratio depends on your application

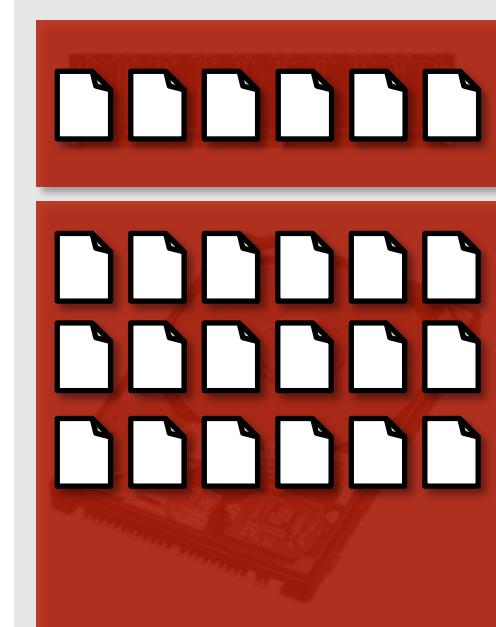
working/total set = .01



COUCHBASE Server

Late stage social game
Many users no longer active; few logged in at any given time.

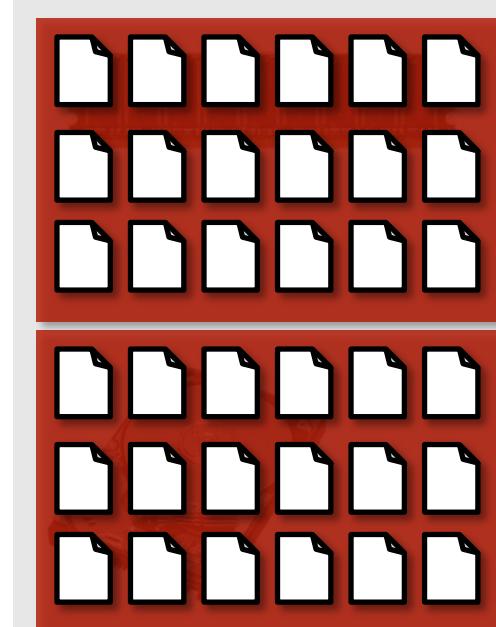
working/total set = .33



COUCHBASE Server

Business application
Users logged in during the day. Day moves around the globe.

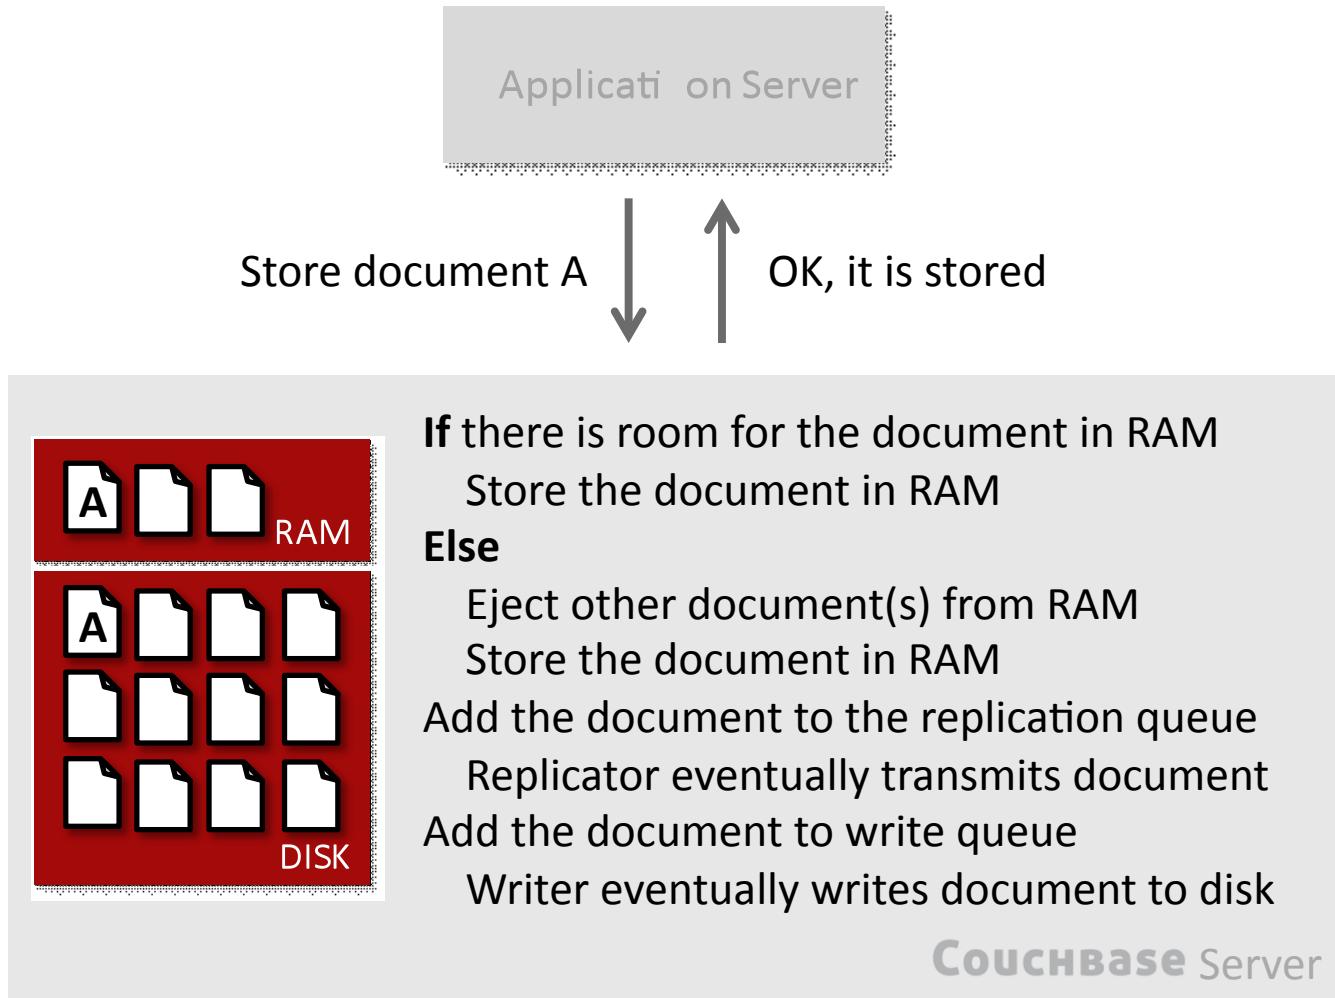
working/total set = 1



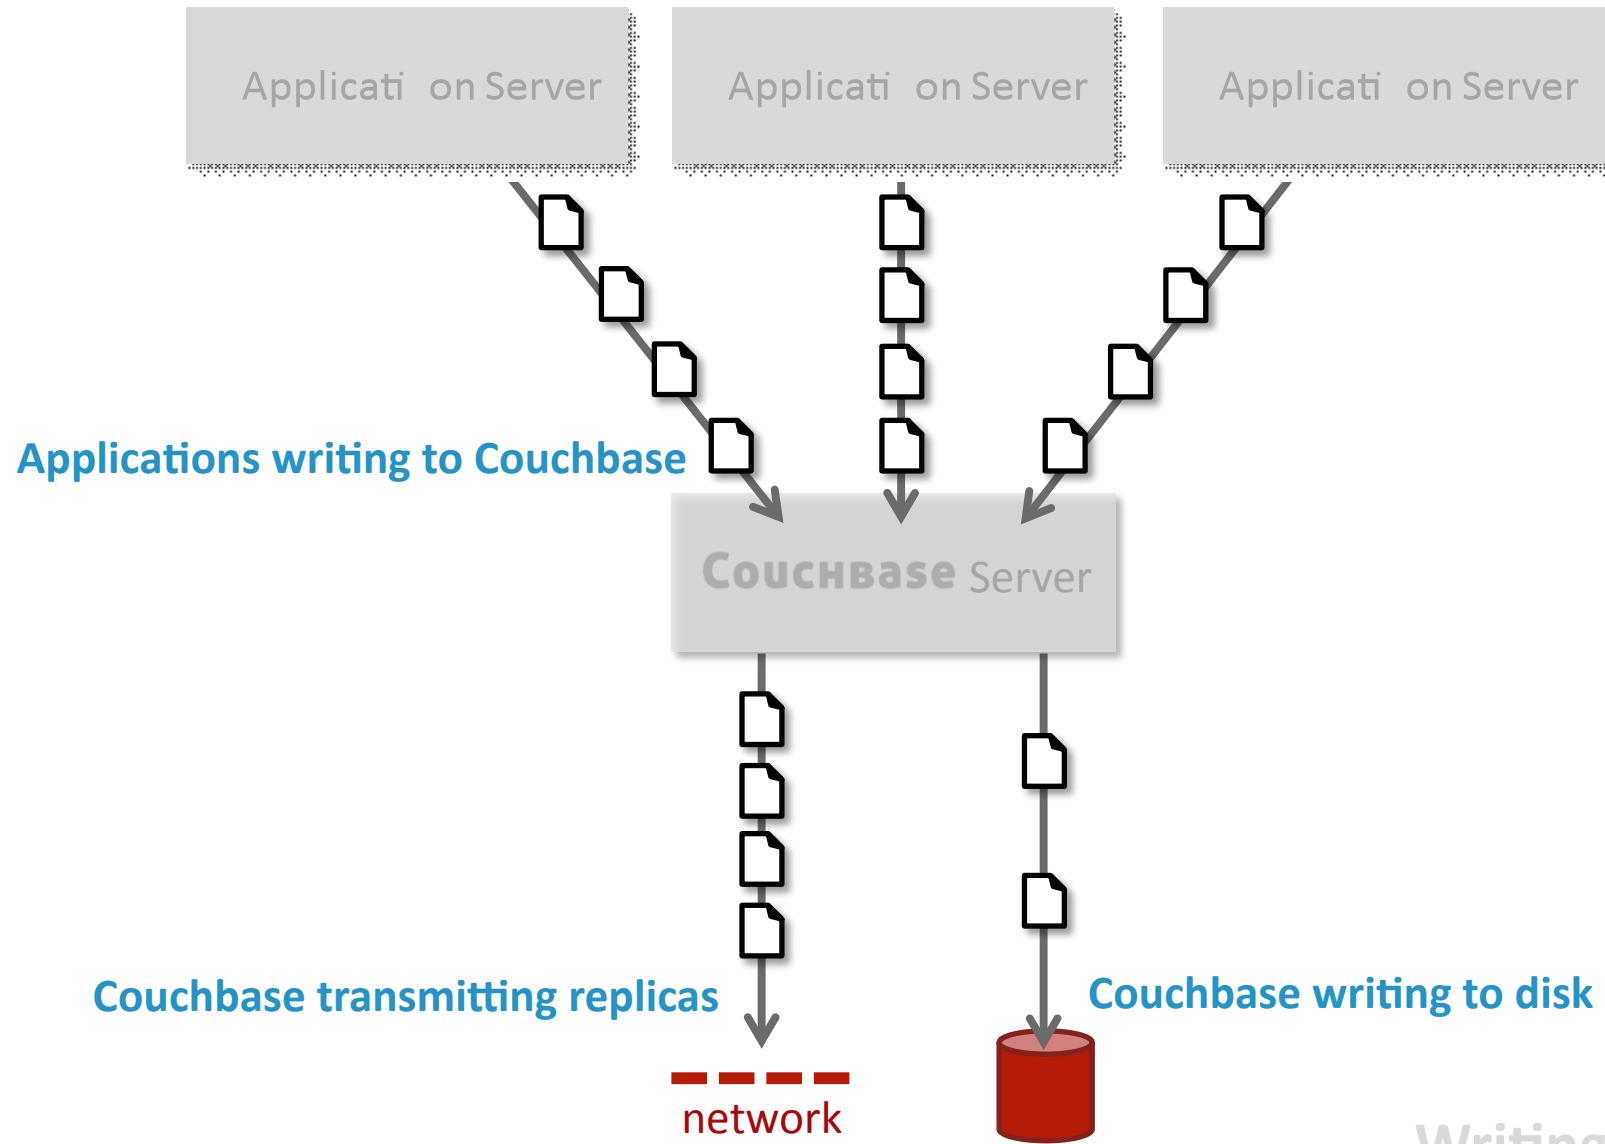
COUCHBASE Server

Ad Network
Any cookie can show up at any time.

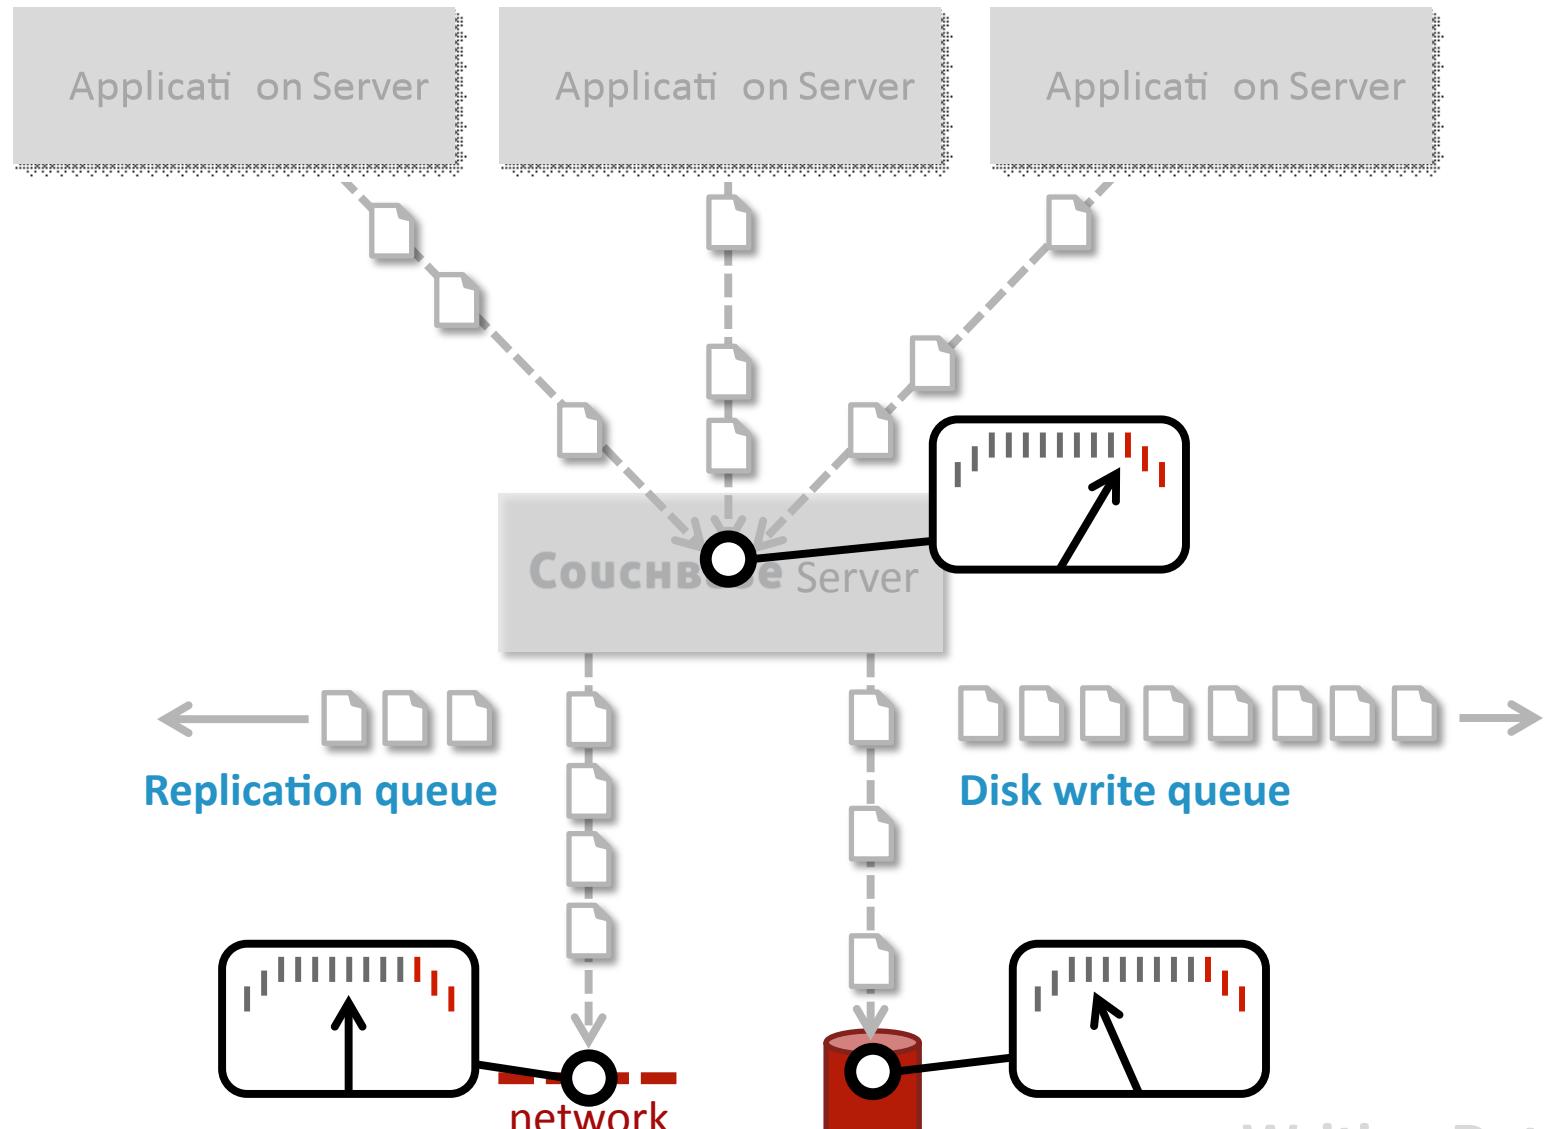
Couchbase in operation: Writing data



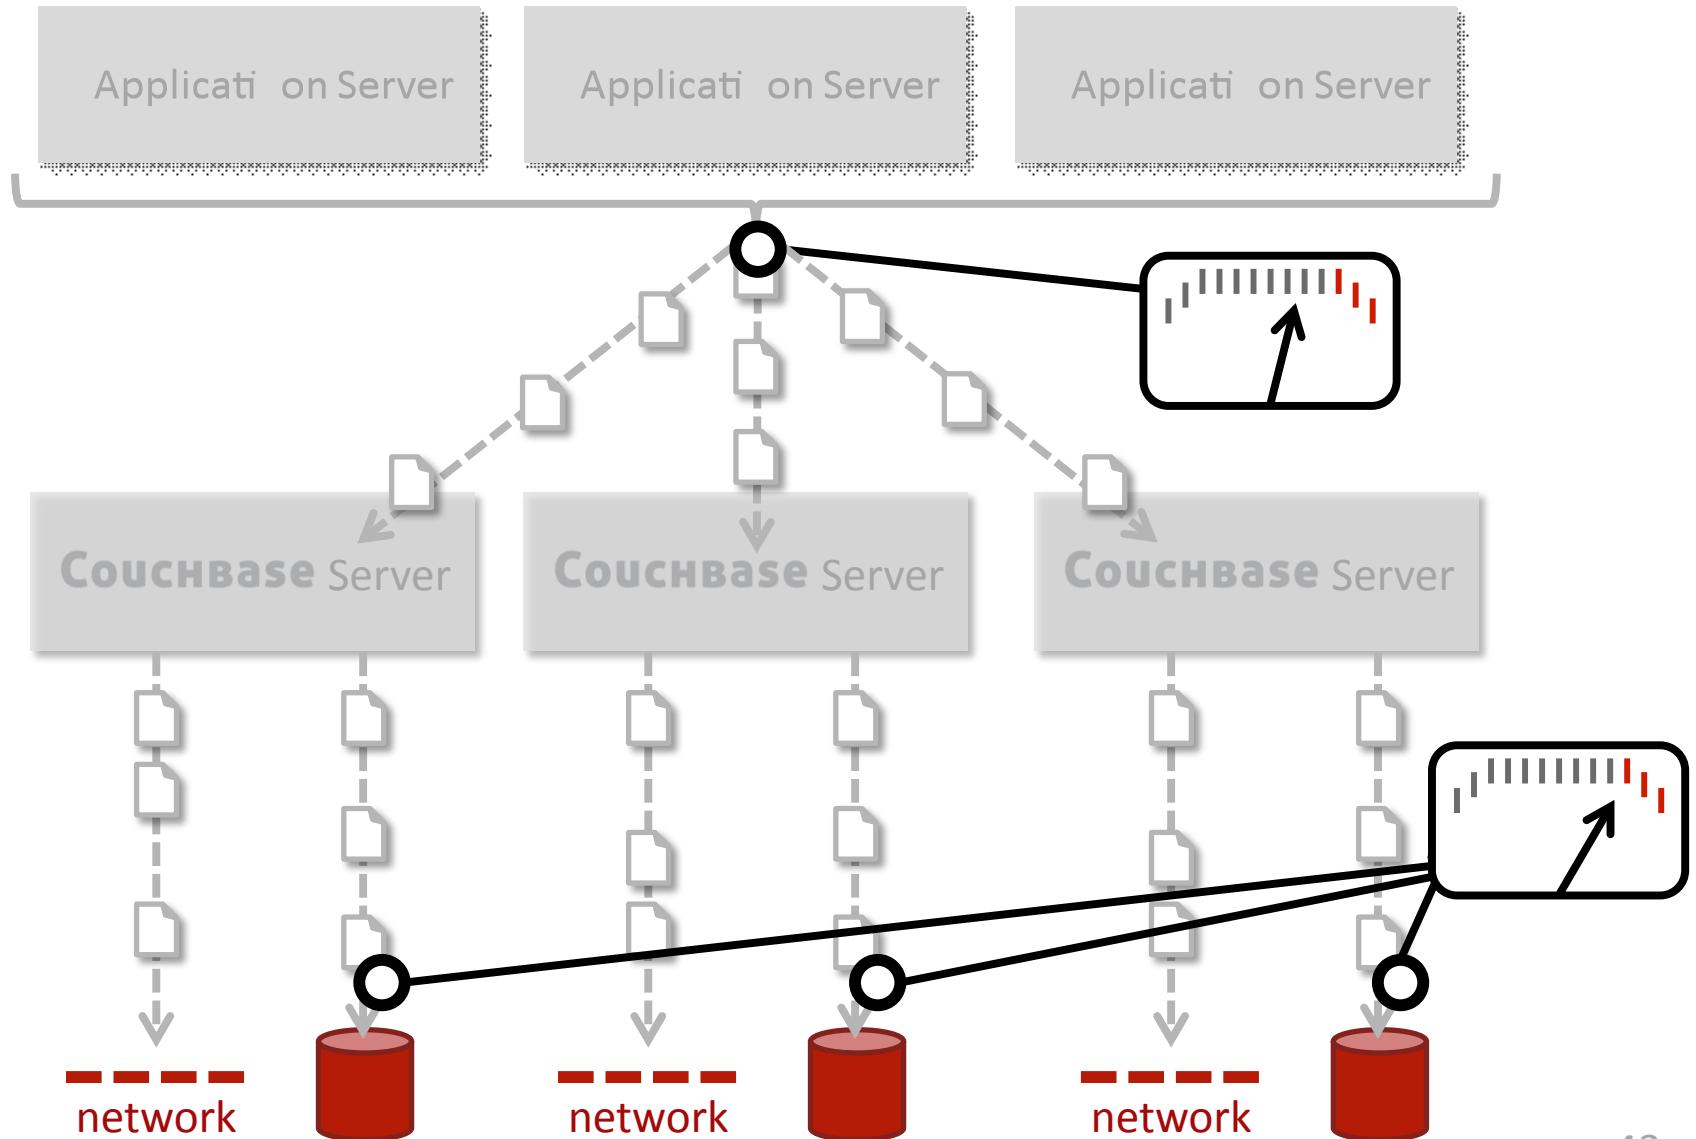
Flow of data when writing



Queues build if aggregate arrival rate exceeds drain rates



Scaling out permits matching of aggregate flow rates so queues do not grow



QUESTIONS?

JAMES@COUCHBASE.COM